

Applying Machine Translation to Two-Stage Cross-Language Information Retrieval

Atsushi Fujii and Tetsuya Ishikawa

University of Library and Information Science
1-2 Kasuga, Tsukuba, 305-8550, Japan

E-mail: fujii@ulib.ac.jp

Abstract. Cross-language information retrieval (CLIR), where queries and documents are in different languages, needs a translation of queries and/or documents, so as to standardize both of them into a common representation. For this purpose, the use of machine translation is an effective approach. However, computational cost is prohibitive in translating large-scale document collections. To resolve this problem, we propose a two-stage CLIR method. First, we translate a given query into the document language, and retrieve a limited number of foreign documents. Second, we machine translate only those documents into the user language, and re-rank them based on the translation result. We also show the effectiveness of our method by way of experiments using Japanese queries and English technical documents.

1 Introduction

The number of machine readable texts accessible via CD-ROMs and the World Wide Web has been rapidly growing. However, since the content of each text is usually provided in a limited number of languages, the notion of information retrieval (IR) has been expanded so that users can retrieve textual information (i.e., documents) across languages. One application, commonly termed “cross-language information retrieval (CLIR)”, is the retrieval task where the user presents queries in one language to retrieve documents in another language. Thus, as can be predicted, CLIR needs to standardize queries and documents into a common representation, so that monolingual IR techniques can be applied. From this point of view, existing CLIR can be classified into three approaches.

The first approach translates queries into the document language [2, 4, 5, 16], while the second approach translates documents into the query language [13, 17]. The third approach projects both queries and documents into a language-independent representation by way of thesaurus classes [6, 18] and latent semantic indexing [3, 11].

Although extensive comparative experiments among different approaches in a rigorous manner are difficult and expensive, a few cases can be found in past CLIR literature.

Oard [17] compared the query and document translation methods. For the purpose of English-German CLIR experiments, he used the 21 English queries

and SDA/NZZ German collection consisting of 251,840 newswire articles, contained in the TREC-6 CLIR collection. Then, he showed that the MT-based query translation with the Logos system was more effective than various types of dictionary-based query translation methods, and that the MT-based document translation method further outperformed the MT-based query translation method. Those findings were salient especially when the length of queries was large.

McCarley [13] conducted English/French bidirectional CLIR experiments, where the 141,656 AP English documents and 212,918 SDA French documents in the TREC-6 and TREC-7 collections were used, and applied a statistical MT method to both query and document translation methods. He showed that the relative superiority between query and document translation methods varied depending on the source and target language pair. To put it more precisely, in his case, the quality of French-English translation was better than that of English-French translation, for both query and document translations.

In addition, he showed that a hybrid method, where the relevance degree of each document (i.e., the “score”) is the mean of those obtained with query and document translation methods, outperformed methods based on either query or document translation, irrespective of the source and target language pair. Possible rationales include that since machine translation is not an invertible operation, query and document translations mutually enhance the possibility that query terms correspond to appropriate translations in documents.

To sum up, the MT-based document translation approach is potentially effective in terms of retrieval accuracy. Besides this, since retrieved documents are mostly in a user’s non-native language, the document translation approach is significantly effective for browsing and interactive retrieval.

However, a major drawback of this approach is that the full translation on large-scale collections is prohibitive in terms of computational cost. In fact, Oard [17], for example, spent approximately ten machine-months in translating the SDA/NZZ collection. This problem is especially crucial in the case where the number of user languages is large, and documents are frequently updated as in the Web. Although a fast MT method [14] was proposed, this method is currently limited to MT within European languages, which are relatively similar to one another.

In view of the above discussions, we propose a method to minimize the computational cost required for the MT-based document translation, which is fundamentally twofold. First, we translate the query into the document language, and retrieve a fixed number of top-ranked documents (one thousand, for example). Second, we machine translate those documents into the query language, and then re-rank those documents based on the score, combining those individually obtained with query and document translation methods. Consequently, it is expected that the retrieval accuracy is improved with a minimal MT cost.

From a different perspective, our method can be classified as a *two-stage* retrieval principle. However, in the monolingual two-stage IR, the second stage usually involves re-calculation of term weights and local feedback so as to increase

the number of relevant documents in the final result [10], and that in the case of existing two-stage CLIR, multiple stages are used to improve the quality of query translation [1, 4].

Section 2 describes our two-stage CLIR system, where we elaborate mainly on the MT-based re-ranking method. Section 3 then evaluates the performance of our system, using the NACSIS test collection [8], which consists of 39 Japanese queries and approximately 330,000 technical abstracts in English and Japanese.

2 System Description

2.1 Overview

Figure 1 depicts the overall design of our Japanese/English bidirectional CLIR system, in which we combined query and document translation modules with a monolingual retrieval system. In this section, we explain the retrieval process based on this figure.

First, given a query in the source language (S), a query translation is performed to output a translation in the target language (T). In this phase, we use two alternative methods. The first method is the use of an MT system, for which we use the Transer Japanese/English MT system.¹ This MT system uses a general bilingual dictionary consisting of 230,000 entries, and 19 optional technical dictionaries, among which a computer terminology dictionary consisting of 100,000 entries is combined with our system.

However, since in most cases, queries consist of a small number of keywords and phrases, word/phrased-based translation methods are expected to be comparable with MT systems, in terms of query translation. Thus, for the second method, we use the Japanese/English phrase-based translation method proposed by Fujii and Ishikawa [5], which uses general/technical dictionaries to derive possible word/phrase translations, and resolves translation ambiguity based on statistical information obtained from the target document collection. In addition, for words unlisted in dictionaries, transliteration is performed to identify phonetic equivalents in the target language.

Second, the monolingual retrieval system searches a collection for documents relevant to the translated query, and sorts them according to the degree of relevance (i.e., the score), in descending order. For English documents, we use the SMART system [19], where the augmented TF·IDF term weighting method (“atc”) is used for both queries and documents, and the score is computed based on the similarity between the query and each document in a term vector space. For Japanese documents, we implemented a retrieval system based on the vector space model.

Consequently, only the top N documents are selected as an intermediate retrieval result, where N is a parametric constant.

Third, the top N documents are translated into the source language. Note that unlike the query translation phase, we use solely the Transer MT system,

¹ Developed by NOVA, Inc.

because translations are aimed primarily at human users, and thus the phrase-based translation method potentially degrades readability of retrieval results.

Finally, the N documents translated are *re-ranked* according to the new score. To accomplish this task, we compute the similarity score between the source query (submitted by the user) and each translated document in the term vector space, as performed in the first retrieval stage. We then compute the new score by averaging those obtained independently with English and Japanese monolingual similarity computations. We will elaborate on this process in Section 2.2.

Note that by decreasing the value of N , we can decrease the computational cost required for machine translation. However, this also decreases the number of relevant documents contained in the top N set, and potentially dilutes the effectiveness of the re-ranking. For example, in an extreme case where the top N set contains no relevant document, the re-ranking procedure does not change the retrieval accuracy.

The re-ranking procedure is similar to McCarley’s hybrid method [13], in the sense that his method also combines scores obtained with query and document translations. However, unlike McCarley’s method, which needs to translate the entire document collection prior to the retrieval, in our method the overhead for translating documents is minimized and can be distributed to each user. In other words, the second stage can be performed on each client (i.e., users’ computers or Web browsers). In fact, there are a number of commercial Web browsers combined with MT systems, and thus it is feasible to additionally introduce the re-ranking function to those browsers. Besides this, we can easily replace the MT system with a newer version or those for other language pairs.

2.2 MT-based Re-ranking Method

First, given the top N documents retrieved and translated into the source language, we first compute the similarity score between each document and the source query provided by the user. Following the vector space model, both queries and documents are represented by a vector consisting of statistical factors associated with indexed terms (i.e., term weights).

In conventional retrieval systems, documents are indexed to produce an inverted file, prior to the retrieval, so that documents containing query terms can efficiently be retrieved even from a large-scale collection. However, in the case of our re-ranking process, since (a) the number of target documents is limited, and (b) real-time indexing degrades the time efficiency, we prefer to use a simple pattern matching method, instead of the inverted file.

For term weighting, we tentatively use a variation of TF-IDF [20, 23], as shown in Equation (1).

$$\begin{aligned} TF &= 1 + \log(f_{t,d}) \\ IDF &= \log \frac{N}{n_t} \end{aligned} \tag{1}$$

Here, $f_{t,d}$ denotes the frequency that term t appears in document d . Note that unlike the common IDF formula, N denotes the number of documents retrieved




Fig. 1. The overall design of our CLIR system.

in the first stage (see Section 2.1), and n_t denotes the number of documents containing term t , out of N documents.

One may argue that since in our case where the number of target documents is considerably smaller than that of the entire collection, a different term weighting method is needed. For example, the IDF formula proposed for large-scale document collections may be less effective for a limited number of documents. However, a preliminary experiment showed that the use of IDF marginally improved the performance obtained without IDF. On the other hand, since the preliminary experiment showed that the use of document length considerably degraded the performance, we compute the similarity between the query and each document, as the inner product (instead of the cosine of the angle) between their associated vectors.

Thereafter, for each document, we combine two similarity scores obtained in English-English and Japanese-Japanese retrieval processes. We shall call them *ESIM* and *JSIM*, respectively. Since those two similarity scores have different ranges, we use a geometric mean, instead of an arithmetic mean, as shown in Equation (2).

$$SIM = ESIM^\alpha \cdot JSIM^\beta \quad (2)$$

Here, *SIM* is the final similarity score with which we re-rank the top N doc-

uments, and α and β are parametric constants used to control the degree to which *ESIM* and *JSIM* affect the computation of *SIM*. However, in the case where either *ESIM* or *JSIM* is zero, the value of *SIM* always becomes zero, disregarding the value of the other similarity score. To avoid this problem, in such a case we arbitrarily assign the value 0.0001 to either *ESIM* or *JSIM* that takes zero.

Possible factors to set values of α and β include the quality of Japanese-English and English-Japanese translations. In the case where the quality of one of the translations is considerably lower, α and β must be properly set so as to decrease the effect of the similarity score through the lower quality translation. Generally speaking, the quality of English-Japanese translation is higher than that of Japanese-English translation, because morphological and syntactic analyses for Japanese are usually more crucial than those for English. However, we empirically set $\alpha = \beta = 1$, that is, we consider *ESIM* and *JSIM* equally in the re-ranking process.

3 Experimentation

3.1 Methodology

We investigated the performance of several versions of our system in terms of Japanese-English CLIR, where each system outputs the top 1,000 documents, and the TREC evaluation software was used to calculate non-interpolated average precision values.

For the purpose of our experiments, we used the official version of the NACSIS test collection [8]. This collection consists of 39 Japanese queries and approximately 330,000 documents (in either a combination of English and Japanese or either of the languages individually), collected from technical papers published by 65 Japanese associations for various fields.

Each document consists of the document ID, title, name(s) of author(s), name/date of conference, hosting organization, abstract and keywords, from which titles, abstracts and keywords were indexed by the SMART system. We used as target documents 187,081 entries that are in both English and Japanese.

Each query consists of the query ID, title of the topic, description, narrative and list of synonyms, from which we used only the description. Figure 2 shows example descriptions (translated into English by one of the authors).

The NACSIS collection was produced for a TREC-type (CL)IR workshop held by NACSIS (National Center for Science Information Systems, Japan) in 1999.² In this workshop, each participant was allowed to submit more than one retrieval result using different methods. However, at least one result had to be gained with only the description field in queries. According to experimental results reported in the proceedings of the workshop [15], in the case where only

² See <http://www.rd.nacsis.ac.jp/~ntcadm/workshop/work-en.html> for details of the NACSIS workshop.

the description field was used, average precision values ranged from 0.021 to 0.182.

Relevance assessment was performed based on the pooling method [22]. To put it more precisely, candidates for relevant documents were first pooled by multiple retrieval systems (primarily systems that participated in the NACSIS workshop). Thereafter, for each candidate document, human expert(s) assigned one of three ranks of relevance, that is, “relevant”, “partially relevant” and “irrelevant”. The average number of candidate documents pooled for each query is 2,509, among which the number of relevant and partially relevant documents are approximately 21 and 6, respectively. In our experiments, we did not regard “partially relevant” documents as relevant ones, because interpretation of “partially relevant” is not fully clear to the authors. Note that since the NACSIS collection does not contain English queries, we cannot estimate a baseline for Japanese-English CLIR performance using English-English IR.

In the following two sections, we will show experimental results in terms of the first and second stages (i.e., query translation methods and the MT-based re-ranking method), respectively.

ID	Description
0032	middleware construction in network collaboration
0035	digital libraries in distributed systems
0036	problems related to groupwares in mobile communication
0062	life-long education and volunteer
0065	image retrieval based on genetic algorithm

Fig. 2. Example query descriptions in the NACSIS collection.

3.2 Evaluation of Query Translation Methods

The primal objective in this section is to compare the effectiveness of the phrase-based translation method proposed by Fujii and Ishikawa [5] and one based on the Transer MT system, in terms of Japanese-English query translation. While the former method is aimed solely at words and phrases, the MT system can also be used for full sentences. In addition, since both methods are, to some extent, complementary to each other, we theoretically gain a query expansion effect, combining query terms translated by individual methods. In view of those above factors, we compared the following query translation methods:

- the use of the Transer MT system for full sentences contained in the description field (“MTS”),
- the use of the Transer MT system for content words and phrases extracted from the description field, for which the ChaSen morphological analyzer [12] was used (“MTP”),

- the phrase-based translation method applied to the same words and phrases as used for the MTP method (“PBT”),
- the use of query terms obtained with both MTP and PBT, where terms outputted by both methods are considered to appear twice in the query (“MPBT”).

Table 1 shows the non-interpolated average precision values, averaged over the 39 queries, for different query translation methods listed above. The second column denotes the average number of query terms provided with each translation method, some of which were potentially discarded as stopwords by the SMART system. The third column denotes average precision values for different query translation methods. We will explain the fourth and fifth columns in Section 3.3.

Looking at this table, one can see that while two MT-based methods, that is, MTS and MTP, were quite comparable in performance, and that PBT outperformed both of them. In the case of PBT, the transliteration successfully identified English equivalents for *katakana* words unlisted in the word dictionary, such as “*coraboreishon* (collaboration)” and “*mobairu* (mobile)”, which the MT-based methods failed to translate. Another reason was due to the difference in dictionaries used. Generally speaking, PBT tended to output technical words more than the MT-based methods. For example, for Japanese phrases “*fukusuudeeta*” and “*sekitsui-doubutsu*”, PBT outputted “multiple data” and “craniate”, while MTS/MTP outputted “more than one data” and “vertebrate”, respectively. Note that this effect was evident partially because the NACSIS collection consists of technical documents. In addition, MPBT further improved the performance of PBT. Although the difference between PBT and MPBT was marginal, it is worth utilizing both the MT-based and phrase-based methods, if available, for query translation.

Table 1. Non-interpolated average precision values, averaged over the 39 queries.

Query Translation			Avg. Precision with Re-ranking	
Method	# of Terms	Avg. Precision	MT	HT
MTS	16.6	0.1124	0.1770 (+57.5%)	0.2297 (+104.3%)
MTP	8.7	0.1134	0.1746 (+54.0%)	0.2217 (+95.5%)
PBT	6.1	0.1403	0.2013 (+43.5%)	0.2295 (+63.6%)
MPBT	13.1	0.1426	0.1986 (+39.3%)	0.2356 (+65.2%)

To validate those above results in a thorough manner, we used the non-parametric Wilcoxon matched-pairs signed-test for statistical testing (at the 5% level), which investigates whether the difference in average precision is meaningful or simply due to chance [7, 9, 21]. We found that differences in average precision values for pairs “MTP versus MTS”, “MPBT versus MTS”, and

“MPBT versus MTP” were significant, although for other pairs, we could not obtain sufficient evidence to conclude a statistical significance. To sum up, we concluded that in query translation, a combination of MT-based and phrase-based translation methods was more effective than a method relying solely on the MT system.

3.3 Evaluation of the MT-based Re-ranking Method

First, we consider Table 1 again, where the fourth column “MT” denotes the average precision values for each query translation method, combined with the MT-based re-ranking method. Throughout our experimentation in this paper, the best average precision value by an automatic method was 0.2013 (i.e., one obtained by PBT combined with the MT-based re-ranking method), which is relatively high, when compared with average precision values reported in the NACSIS workshop (ranging from 0.021 to 0.182).

For each query translation method, the improvement in average precision from one without the re-ranking, which is generally noticeable, is indicated in parentheses. In fact, we used the Wilcoxon test again, as conducted in Section 3.2, and confirmed that every improvement was statistically significant. To sum up, the MT-based re-ranking method we proposed was generally effective, irrespective of the query translation method combined, in terms of CLIR performance.

Second, we conducted an error analysis for queries for which the re-ranking method degraded the average precision, and found that roughly two thirds of errors were due to ambiguity in the document translation. For example, the English word “library” was often incorrectly translated into “*raiburari* (library as a software)”, whereas the original query was intended to “*toshokan* (library as an institution)”.

Third, to estimate the upper bound of the re-ranking method, as denoted in the fifth column “HT”, we used as human translations Japanese documents comparable to English ones in the NACSIS collection. By comparing the results of “MT” and “HT”, one can see that MT systems with a higher quality, if available, are expected to further improve our CLIR system. In fact, when we manually corrected inappropriate translations in translated documents, such as “library (*raiburari/toshokan*)” above, the average precision of “MT” became almost equivalent to that of “HT”.

Noted that when combined with the re-ranking method, differences among query translation methods in average precision were relatively overshadowed. In the case of “MT”, the Wilcoxon test showed that differences in only pairs “MPBT versus MTS” and “MPBT versus MTP” were significant, while in the case of “HT”, none of the differences were identified as significant.

Fourth, we investigated how the number of documents retrieved in the first stage (i.e., the value of N in Section 2) affected the performance of the re-ranking method. As discussed in Section 2.1, in real world usage, one has to consider the trade-off between the retrieval accuracy (i.e., average precision in our case) and overhead required for the document translation.

Table 2 shows the results, where average precision values in the column “1,000” correspond to those in Table 1. By comparing average precision values for each of four query translation methods (i.e., MTS, MTP, PBT and MPBT) and those suffixed with “+MT” and “+HT” in Table 2, one can see that the re-ranking methods were effective, irrespective of the number of documents retrieved. In other words, it is expected that we can minimize the overhead in translating documents, without decreasing the retrieval accuracy.

Table 3 shows CPU time (sec.) required for the document translation and re-ranking procedures, averaged over four different query translation methods. In the case of $N = 1,000$, the total CPU time was approximately three minutes, which is perhaps not tolerable for a real-time usage. However, for small values of N (e.g., 50 and 100), the CPU time was more acceptable and practical, maintaining the improvement of retrieval accuracy.

Table 2. The relation between the number of documents retrieved in the first stage and non-interpolated average precision values, averaged over the 39 queries.

Method	# of Documents Retrieved (N)						
	50	100	200	400	600	800	1,000
MTS	0.0949	0.1017	0.1074	0.1101	0.1112	0.1119	0.1124
MTS+MT	0.1341	0.1556	0.1673	0.1698	0.1720	0.1736	0.1770
MTS+HT	0.1666	0.1901	0.2070	0.2173	0.2230	0.2259	0.2297
MTP	0.0953	0.1020	0.1085	0.1113	0.1123	0.1131	0.1134
MTP+MT	0.1449	0.1584	0.1692	0.1711	0.1728	0.1750	0.1746
MTP+HT	0.1619	0.1819	0.2017	0.2105	0.2165	0.2203	0.2217
PBT	0.1215	0.1301	0.1355	0.1385	0.1394	0.1399	0.1403
PBT+MT	0.1553	0.1723	0.1866	0.1954	0.1978	0.2005	0.2013
PBT+HT	0.1722	0.1915	0.2097	0.2212	0.2241	0.2279	0.2295
MPBT	0.1229	0.1305	0.1376	0.1405	0.1416	0.1421	0.1426
MPBT+MT	0.1690	0.1766	0.1901	0.1946	0.1958	0.1967	0.1986
MPBT+HT	0.1814	0.1968	0.2142	0.2242	0.2301	0.2319	0.2356

4 Conclusion

Reflecting the rapid growth in utilization of machine readable texts, cross-language information retrieval (CLIR) has variously been explored in order to facilitate retrieving information across languages.

In brief, existing CLIR systems are classified into three approaches: (a) translating queries into the document language, (b) translating documents into the query language, and (c) representing both queries and documents in a language-independent space. Among these approaches, the second approach, based on machine translation, is effective in terms of retrieval accuracy and user interaction.

Table 3. CPU time for document translation and re-ranking (sec.).

	# of Documents Retrieved (N)						
	50	100	200	400	600	800	1,000
translation	9.5	17.7	33.3	65.6	106.2	139.3	175.1
re-ranking	0.2	0.3	0.6	1.2	1.8	2.4	3.0
total	9.7	18.0	33.9	66.8	108.0	141.7	178.1

(Pentium III 700MHz)

However, the computational cost in translating large-scale document collections is prohibitive.

To resolve this problem, we proposed a two-stage CLIR method, in which we first used a query translation method to retrieve a fixed number of documents, and then applied machine translation only to those documents, instead of the entire collection, to improve the document ranking.

Through Japanese-English CLIR experiments using the NACSIS collection, we showed that our two-stage method significantly improved average precision values obtained solely with query translation methods. We also showed that our method performed reasonably, even in the case where the number of retrieved documents was relatively small.

Acknowledgments

The authors would like to thank NOVA, Inc. for their support with the Transer MT system, and Noriko Kando (National Institute of Informatics, Japan) for her support with the NACSIS collection.

References

- [1] Lisa Ballesteros and W. Bruce Croft. Phrasal translation and query expansion techniques for cross-language information retrieval. In *Proceedings of the 20th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 84–91, 1997.
- [2] Lisa Ballesteros and W. Bruce Croft. Resolving ambiguity for cross-language retrieval. In *Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 64–71, 1998.
- [3] Jaime G. Carbonell, Yiming Yang, Robert E. Frederking, Ralf D. Brown, Yibing Geng, and Danny Lee. Translingual information retrieval: A comparative evaluation. In *Proceedings of the 15th International Joint Conference on Artificial Intelligence*, pp. 708–714, 1997.
- [4] Mark W. Davis and William C. Ogden. QUILT: Implementing a large-scale cross-language text retrieval system. In *Proceedings of the 20th Annual International*

ACM SIGIR Conference on Research and Development in Information Retrieval, pp. 92–98, 1997.

- [5] Atsushi Fujii and Tetsuya Ishikawa. Cross-language information retrieval for technical documents. In *Proceedings of the Joint ACL SIGDAT Conference on Empirical Methods in Natural Language Processing and Very Large Corpora*, pp. 29–37, 1999.
- [6] Julio Gonzalo, Felisa Verdejo, Carol Peters, and Nicoletta Calzolari. Applying EuroWordNet to cross-language text retrieval. *Computers and the Humanities*, Vol. 32, pp. 185–207, 1998.
- [7] David Hull. Using statistical testing in the evaluation of retrieval experiments. In *Proceedings of the 16th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 329–338, 1993.
- [8] Noriko Kando, Kazuko Kuriyama, and Toshihiko Nozue. NACSIS test collection workshop (NTCIR-1). In *Proceedings of the 22nd Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 299–300, 1999.
- [9] E. Michael Keen. Presenting results of experimental retrieval comparisons. *Information Processing & Management*, Vol. 28, No. 4, pp. 491–502, 1992.
- [10] K.L. Kwok and M. Chan. Improving two-stage ad-hoc retrieval for short queries. In *Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 250–256, 1998.
- [11] Michael L. Littman, Susan T. Dumais, and Thomas K. Landauer. Automatic cross-language information retrieval using latent semantic indexing. In Gregory Grefenstette, editor, *Cross-Language Information Retrieval*, chapter 5, pp. 51–62. Kluwer Academic Publishers, 1998.
- [12] Yuji Matsumoto, Akira Kitauchi, Tatsuo Yamashita, Osamu Imaichi, and Tomoaki Imamura. Japanese morphological analysis system ChaSen manual. Technical Report NAIST-IS-TR97007, NAIST, 1997. (In Japanese).
- [13] J. Scott McCarley. Should we translate the documents or the queries in cross-language information retrieval? In *Proceedings of the 37th Annual Meeting of the Association for Computational Linguistics*, pp. 208–214, 1999.
- [14] J. Scott McCarley and Salim Roukos. Fast document translation for cross-language information retrieval. In *Proceedings of the 3rd Conference of the Association for Machine Translation in the Americas*, pp. 150–157, 1998.
- [15] National Center for Science Information Systems. *Proceedings of the 1st NTCIR Workshop on Research in Japanese Text Retrieval and Term Recognition*, 1999.
- [16] Jian-Yun Nie, Michel Simard, Pierre Isabelle, and Richard Durand. Cross-language information retrieval based on parallel texts and automatic mining of parallel texts from the Web. In *Proceedings of the 22nd Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 74–81, 1999.
- [17] Douglas W. Oard. A comparative study of query and document translation for cross-language information retrieval. In *Proceedings of the 3rd Conference of the Association for Machine Translation in the Americas*, pp. 472–483, 1998.
- [18] Gerard Salton. Automatic processing of foreign language documents. *Journal of the American Society for Information Science*, Vol. 21, No. 3, pp. 187–194, 1970.
- [19] Gerard Salton. *The SMART Retrieval System: Experiments in Automatic Document Processing*. Prentice-Hall, 1971.
- [20] Gerard Salton and Christopher Buckley. Term-weighting approaches in automatic text retrieval. *Information Processing & Management*, Vol. 24, No. 5, pp. 513–523,

- 1988.
- [21] Padmini Srinivasan. A comparison of two-poisson, inverse document frequency and discrimination value models of document representation. *Information Processing & Management*, Vol. 26, No. 2, pp. 269–278, 1990.
 - [22] Ellen M. Voorhees. Variations in relevance judgments and the measurement of retrieval effectiveness. In *Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 315–323, 1998.
 - [23] Justin Zobel and Alistair Moffat. Exploring the similarity space. *ACM SIGIR FORUM*, Vol. 32, No. 1, pp. 18–34, 1998.