Extended Abstract – Model-Based Debugging of Java Programs[*][†]

Cristinel Mateis   Markus Stumptner   Dominik Wieland

Franz Wotawa [‡]

Technische Universität Wien

Institut für Informationssysteme

Database and Artificial Intelligence Group

Favoritenstraße 9–11, A-1040 Vienna, Austria

{mateis,mst,wieland,wotawa}@dbai.tuwien.ac.at

http://www.dbai.tuwien.ac.at

1   Introduction

Model-based reasoning is a central concept in current research into intelligent diagnostic systems. It is based on the assumption that sources of incorrect behavior in technical devices can be located and identified via the existence of a model describing the basic properties of components of a certain application domain. When actual data concerning the misbehavior of a system composed from such components is available, a domain-independent diagnosis engine can be used to infer which parts of the system contribute to the observed behavior. Model-based Diagnosis provides a set of proven algorithms and methods for searching faults [14] and identifying points of measurement [3].

[‡]Authors are listed in alphabetical order.

This paper describes the application of the model-based approach to the debugging of Java programs written in a subset of Java. We show how a simple dependency model can be derived from a program, demonstrate the use of the model for debugging and reducing the required user interactions, give a comparison of the functional dependency model with program slicing [16], and finally discuss some current research issues.

2   Model-Based Diagnosis

The model-based approach is based on the notion of providing a representation of the correct behavior of a technical system. By describing the structure of a system and the function of its components, it is possible to ask for the reasons why the desired behavior was not achieved. In the diagnosis community, the model-based approach has achieved wide recognition due to its advantages:

- once an adequate model has been developed for a particular domain, it can be used to diagnose different actual systems of that domain

- the model can be used to search for single or multiple faults in the system without alteration

- different diagnosis algorithms can be used for a given model

- the existence of a clear formal basis for judging and computing diagnoses

Using the standard consistency-based view as defined by Reiter [14], a diagnosis system can be seen formally as a tuple $(SD, COMP)$ where SD is a logical theory sentence modeling the behavior of the given system (in our case the program to be debugged), and $COMP$ a set of components, i.e., statements. A diagnosis system together with a set of observations OBS, i.e., a test-case, forms a diagnosis problem. A diagnosis Δ, i.e., a bug candidate, is a subset of $COMP$, with the property that the assumption that all statements in Δ are incorrect, and the rest of the statements is correct, should be consistent with SD and OBS. Formally, Δ is a diagnosis iff $SD \cup OBS \cup \{\neg AB(C) | C \in COMP \setminus \Delta\} \cup \{AB(C) | C \in \Delta\}$ is consistent. A component not working as expected, i.e., a statement containing a bug, is represented by the predicate $AB(C)$.

The basis for this is that an incorrect output value (where the incorrectness can be observed directly or derived from observations of other signals) cannot be produced by a correctly functioning component with correct inputs. Therefore, to make a system with observed incorrect behavior consistent with the description and avoid a contradiction, some subset of its components must be assumed to work incorrectly. In practical terms, one is interested in finding minimal diagnoses, i.e., a minimal set of components whose malfunction explains the misbehavior of the system (otherwise, one could explain every error by simply assuming every component to be malfunctioning). Basic properties of the approach as well as algorithms for efficient computation of diagnoses are described in [14].

Starting from straightforward work that used a logic program directly as system descriptions [2, 1], in the last years the use of model-based reasoning (MBR) for debugging of software has been examined in a wider context [12, 4, 15, 10, 11]. All of the approaches have in common that they use a model derived from a program for locating (or, rarely, correcting) a bug. They differ in the considered programming language (ranging from purely logical languages to hardware description languages – in particular VHDL [4], functional, and finally imperative languages), and the type of model (qualitative [12], dependency-, or value-based models). The purpose of previous research was to show the applicability of MBR in the software domain by introducing models, often for special purpose languages. Our current work deals with the extension and application of these principles to a mainstream language (Java). This paper presents first results of an implemented debugger prototype using different example programs. The JADE debugger currently implements a functional dependency model that extends our earlier work [10, 11]. The granularity of the debugger, i.e., the elements of a program that are considered to be faulty or not, is currently set to the statement level (instead of individual expression level) for efficiency reasons.

The JADE debugger combines the standard operation modes of diagnosis systems and standard debuggers. First, the program is converted into the dependency representation which is compiled into a logical model. Once the program has been executed actual observations of its behavior can be provided. This behavior together with the logical model is used by the diagnosis engine to compute bug candidates that map back to positions (and statements) within the program to be debugged (see Figure 1).
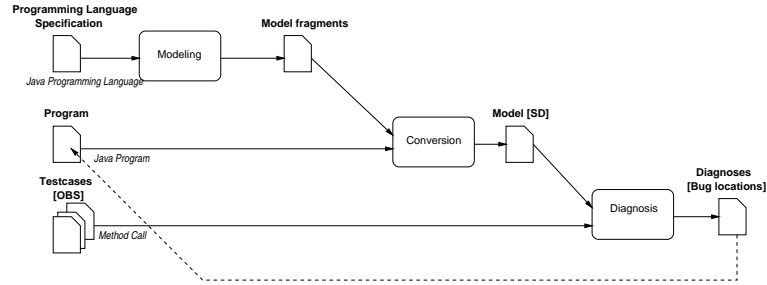
Figure 1: Model-based Diagnosis for Debugging

3   Modeling for Debugging

We first give an overview of computing the dependencies for Java programs and then describe the derivation of the system description which is used for diagnosis.

Dependencies   For the sake of brevity we omit the discussion of object-oriented features such as dynamic binding in this paper and concentrate on the "basic" imperative features of the language. A more detailed technical description of the basic idea behind the conversion algorithm (excluding external side effects and method calls) can be found in [10], and a discussion about references and side effects in [11]. In order to compute dependencies, we must consider that the variables occurring in the methods change their values during program execution. This is handled by assigning a unique index to all locations where a variable occurs as target of an assignment (we refer to each such location as an *occurrence* of the variable. It is the various variable occurrences that dependencies are computed for.

Let x and y be indexed variable occurrences of a given method m. We say that x *depends on* y iff the value of y determines the value of x for at least one input vector. This definition is based on earlier work on dependencies in software debugging, e.g., [8, 13]. Beside debugging dependencies are used for verification (see [7]). Formally, we define a functional dependency as a pair (x, M_x), where x is a variable occurrence and M_x is a set of variable occurrences such that x depends on every $y \in M_x$. We can now compute all functional dependencies of a particular statement by determining the functional dependencies of all variables used within the statement.

```
1. class SWExamples {
2.    public static void test(int a,b,c,d,e) {
3.       int f,g,s1,s2,s3;
4.       s1=a*c;
5.       s2=b*d;
6.       s3=c*e;
7.       f=s1+s2;
8.       g=s2+s3;
9.    }
10}
```

(a) Source code

Line	Environment
	test(3,2,2,3,3)
2.	$a_{test} = 3$, $b_{test} = 2$, $c_{test} = 2$, $d_{test} = 3$, $e_{test} = 3$
3.	
4.	$s1_{test} = 6$
5.	$s2_{test} = 6$
6.	$s3_{test} = 6$
7.	$f_{test} = 12$
8.	$g_{test} = 12$
9.	
10.	$test(3, 2, 2, 3, 3) = void$

(b) Evaluation Trace for
test(3,2,2,3,3)

Figure 2: A simple Java method

Whereas the variables correspond to the ports of traditional diagnosis components, the natural choice for components given the abstract nature of the dependency-based representation are statements. The set of diagnosis components can be viewed as a diagnosis system where the connections are formed by the variable occurrences inside the components, i.e., components c_i and c_j are connected iff one component establishes the functional dependency (v_i, M) and the other, $(w_j, \{\ldots, v_i, \ldots\})$. Variable occurrences v_0 are inputs of the whole diagnosis system. A variable occurrence v_i is an output iff there is no other occurrence v_j such that $j > i$. Since during conversion indices are always increased, the resulting diagnosis system, i.e., the graph representation, is acyclic.

Computing functional dependencies for Java programs requires compiling each method declared for a class. A method m of a Java program is converted by sequentially converting its statements into diagnosis components.

We illustrate the computation of functional dependencies using the example of Figure 2. The functional dependencies for `s1=a*c;` are $\{(s1_2, \{a_1, c_1\})\}$ because the value of `s1` is given by the product of the values of `a` and `c`. In summary we obtain the following dependency sets (indices are ignored) for the 5 statements: $fd(C_4) = \{(s1, \{a, c\})\}$, $fd(C_5) = \{(s2, \{b, d\})\}$, $fd(C_6) = \{(s3, \{c, e\})\}$, $fd(C_7) = \{(f, \{s1, s2\})\}$, $fd(C_8) = \{(g, \{s2, s3\})\}$

where C_i denotes the statement in line i.

The System Description After computing all dependencies, we map them to a logical representation, which can be directly used for model-based debugging. For this purpose we assume that the statements are given as a set $COMP$, and that for all statements, the functional dependencies are defined. The set of functional dependencies for a statement st is written as $fd(st)$.

Functional dependencies describe behavior implicitly by describing influences between variables. Instead of speaking about real values, we can only speak about whether a value v is correct (written as $ok(v)$) or not (written $nok(v)$). We can further write that if a statement s is assumed to be correct (i.e., $\neg AB(s)$ holds) and all input variables have a correct value then the value of variables used as target in an assignment statement must be correct. Formally, the system description is given by:

$$\forall_{(o,M)\in fd(C)} \left[\neg AB(C) \wedge \bigwedge_{x\in M} ok(x) \rightarrow ok(o) \right] \in SD$$

where $C \in COMP$ is a statement. In addition, we know that it is impossible that a variable value is known to be correct and incorrect at the same time. Therefore, we have to add the rule $ok(v) \wedge nok(v) \rightarrow \bot$ to the model SD, for each variable occurrence v in the program. The described model can be used together with a standard MBD algorithm for computing bug locations.

For software debugging, the observations required for diagnosis are given by the specified behavior, in our case the expected input/output vectors. By comparing the specified output with the computed output, we can classify the correctness of variables. Variables v that are assumed to have the correct value lead to the observation $ok(v)$. Variables with an incorrect value are represented by $nok(v)$.

In Figure 2(b) the evaluation trace for the call *test(3,2,2,3,3)* for the Java program given in Figure 2(a) is given. The trace only presents the lines of code which are involved in the current evaluation, and the new environments created. To distinguish different local variables they are indexed with the name of the method where they are declared. In this case there is no return value. From the dependencies computed above, we get the logical model SD:

$$\neg AB(C_4) \wedge ok(a) \wedge ok(c) \rightarrow ok(s1) \qquad \neg AB(C_5) \wedge ok(b) \wedge ok(d) \rightarrow ok(s2)$$
$$\neg AB(C_6) \wedge ok(c) \wedge ok(e) \rightarrow ok(s3) \qquad \neg AB(C_7) \wedge ok(s1) \wedge ok(s2) \rightarrow ok(f)$$
$$\neg AB(C_8) \wedge ok(s2) \wedge ok(s3) \rightarrow ok(g)$$

$$ok(a) \wedge nok(a) \rightarrow \bot \quad ok(b) \wedge nok(b) \rightarrow \bot \quad ... \quad ok(s3) \wedge nok(s3) \rightarrow \bot$$

In this example we assume that the method call *test(3,2,2,3,3)* should lead to values *f=12* and *g=0*, i.e., that line 8 should be *g=s2-s3* instead of *g=s2+s3*. For this case, we get observations OBS:

$$ok(a) \wedge ok(b) \wedge ok(c) \wedge ok(d) \wedge ok(e) \wedge ok(f) \wedge nok(g)$$

Using $SD \cup OBS$ we get 3 diagnoses, each pinpointing a single possible bug location: $\{C_5\}, \{C_6\}, \{C_8\}$. The other statements can be ignored in this case. Using the measurement selection algorithm from [3] we can compute the optimal next question to be presented to the user in order to distinguish between the 3 candidates.

4 Diagnosing with the Dependency Model

The JADE debugger with dependency model can be proven to be complete with regard to bugs that do not alter the dependency structure of the program, since all statements that may cause a wrong value are considered and therefore are diagnosis candidates. However, discrimination capability can be low. Consider the example program from Figure 2, together with the specified values $f = 12$ and $g = 0$. In this case, the debugger returns the candidate $\{C5\}$ which could be eliminated when using a value-based model. Now assume $C5$ is incorrect, all other statements are correct, and apply the test case from Figure 2. From $C4$ and $C6$ and the input values we derive $s1_{test} = s3_{test} = 6$. Using these values together with $C7$ and $f_{test} = 12$, we get $s2_{test} = 6$. Now using this value together with the assumption $C8$ is correct leads to $g_{test} = 12$, contradicting our specified value $g_{test} = 0$. Hence, $\{C5\}$ is not a diagnosis w.r.t the value-based model, illustrating the (unsurprising) fact that a model based purely on dependencies is too weak to discriminate between all possible program errors.

Concerning performance, for smaller programs and interactive debugger use, diagnosis times should be in the single second range although longer times are acceptable for very large programs. It is obvious that searching for all single bugs using our model is restricted to $O(n^2)$, where n denotes the number of diagnosis components, i.e., in our case, statements. Using empirical results from [6] we can expect that computing all single bugs for Java methods with several hundred statements should be done in less than

1 second, a result that is consistent with the experience from the VHDL domain [4].

Like program slicing [16], our dependency model is based on static analysis of the code, i.e., it is computed using the program structure and does not use the runtime program behavior for fault localization. A program slice is defined as the part of a program possibly influencing the value of given variables and not occurring within the program after a given position. The slice for our running example for variables $\{g\}$ and position 8 comprises the lines 8, 6, and 5. This result is equal to the one obtained by our dependency model and the question arises about the differences between both approaches. Our dependency model is more hierarchically organized, e.g., formally a conditional statement is viewed as a single diagnosis component and sub-divided only after being identified as faulty. To allow a comparison of slicing with MBD using the functional dependency model, we assume an appropriate mapping. In this respect we obtain similar results from both techniques except in the cases where several variables have a faulty value after program execution. In this situation the model-based approach tries to minimize the source of the misbehavior leading to fewer solutions, while slicing does not. Given the example, assume that line 5 is faulty, leading to wrong values for variables f and g. The slice for f and g is the whole program, while our dependency model would deliver only line 5 as single fault candidate.

5   Empirical results

The following experiments show the results of the JADE debugger using various Java methods from our example library, modified at randomly selected statements:

Example 1: The *adder* method implements a binary full adder mapping three inputs to two outputs.

Example 2: The *library* method is part of a small application that creates a sample library and then computes the author who has published the most books of all authors whose books can be found in the library (methods involving object-oriented language structures, such as multiple objects, instance method calls, class & instance variables, etc...)

Example 3: Sorting methods: various sorting methods, providing the full complement of control statements: loops, selection statements, and method calls.

Test	Method	Lines	Error	Interactions						
				Setup	Query	Loop	Exprs.	Iter.	Total	Total 2
1	*adder_f*1	17	4	1	2	0	1	0	4	3
2	*adder_f*2	17	4	1	2	0	1	0	4	3
3	*adder_f*3	17	7	1	2	0	1	0	4	3
4	*adder_f*4	17	7	1	2	0	1	0	4	3
5	*adder_f*5	17	7	1	2	0	1	0	4	3
6	*adder_f*6	17	12	1	4	0	1	0	6	5
7	*adder_f*7	17	12	1	4	0	1	0	6	5
8	*adder_f*8	17	11	1	4	0	1	0	6	5
9	*adder_f*9	17	11	1	6	0	1	0	8	7
10	*adder_f*10	17	14	1	4	0	1	0	6	5
11	*adder_f*11	17	14	1	4	0	1	0	6	5
12	*adder_f*12	17	16	1	4	0	1	0	6	5
13	*adder_f*13	17	12	1	4	0	1	0	6	5
14	*adder_f*14	17	9	1	4	0	1	0	6	5
15	*library_f*1	30	28	1	5	2	1	1	10	6
16	*bubbleSort_f*1	10	4	1	2	2	1	1	7	3
17	*bubbleSort_f*2	10	7	1	5	3	1	2	12	6
18	*insertionSort_f*	19	4	1	2	2	1	1	7	3
19	*shellSort_f*1	14	3	1	2	1	1	1	6	3
20	*shellSort_f*2	14	6	1	3	3	1	2	10	4
21	*selectionSort_f*1	14	1	1	1	0	1	0	3	2
22	*heapSort_f*1	11	7	1	6	1	1	1	10	7
23	*heapSort_f*2	11	8	1	4	1	1	1	8	5
24	*heapSort_f*3	11	9	1	3	1	1	1	7	4
Sum		382	217	24	81	16	24	11	156	105
Av.		15.92	*9.042*	1	3.37	0.67	1	0.46	6.5	*4.37*

Table 1: Debugging results from Java examples

Table 1 shows (from left to right) the tested method (in which a single error has been installed), the total number of statements in each method, the index of the buggy statement within the method (which can directly be used to compare the outcomes of the Jade debugger tests with "manual" use of a debugger where the user steps through the code sequentially until the erroneous line is found), and finally the number of user interactions which are needed to exactly locate the bug, classified by type:

Setup: verify the system output connections, i.e. all output variables of the method.

Query: verify the value of a variable at a particular statement

Loop: debug the condition of a loop or selection statement

Exprs: find the smallest sub-expression of a particular statement. This allows further debugging of method and constructor calls.

Iter.: determine the first loop iteration in which the error occurs

The "Total" column from the right shows the total number of user inter-

actions, and the "Total2" column the sum of all setup and query interactions. The latter figure determines the debugger's performance at statement level and can therefore be compared with Column 4. These two columns give a simple comparison of the user interactions needed to find the exact location of a faulty statement with the two different debugging strategies, i.e. model-based vs. unsupported debugging.

Locating a bug in the *adder* method with a traditional debugging tool requires 10 user interactions versus 4.43 with the JADE debugger. This fairly drastic difference in favor of the model-based technique is due to the simple block structure of the method with the model-based debugger exploiting its knowledge about the underlying functional dependency structure of the block. The longer the block, the greater the advantage. For the same reason, the faulty version of the *library* method requires 28 user interactions in the traditional approach compared to 6 using the JADE debugger. To locate a bug in a sorting algorithm a traditional debugger on average needs 5.45 steps versus 4.12 for the JADE debugger, the advantage being less clear due to the complex control structures.

Overall, 24 tests are compiled in the table with 9.04 user interactions on average required in the traditional case and 4.37 for the JADE debugger, clearly highlighting the potential abilities of model-based debugging tools.

Note that as with techniques like program slicing, the bug location efficiency of the model-based approach does not depend on the actual location of the error in the code. The number of user interactions it takes to find the bug is therefore much better assessable.

6   Discussion and Conclusion

One of the main advantages of the model-based approach is the ability to incorporate multiple models in the same formal and computational framework. The dependency-based representation described in this paper has the advantage of being simple and computationally efficient to solve, thus providing an answer to the important question of scalability of the approach (our use of a dependency based mechanism in the VHDL domain was usable for very large programs [4]). Its disadvantage of low discrimination can be counteracted by combining it with a more detailed (i.e., value-based) model, e.g., by focusing in on smaller program parts isolated by the dependency model, and diagnosing them with the more detailed model, which is cur-
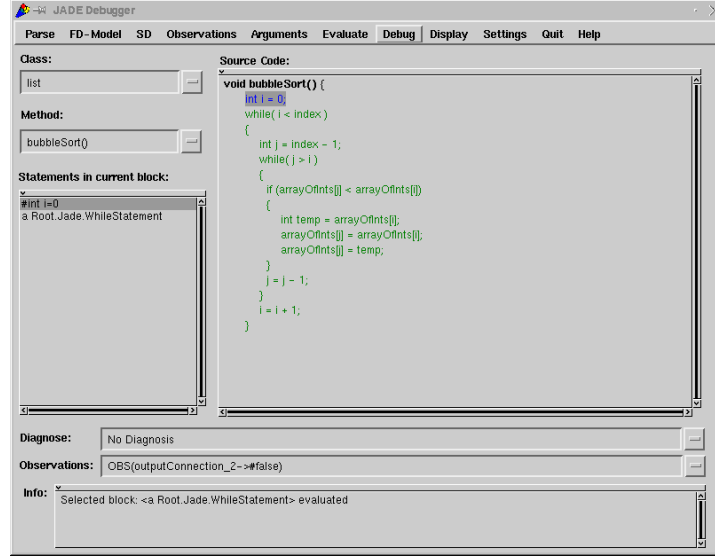
Figure 3: The JADE debugger main window

rently undergoing implementation [11]. The value-based model, unlike the dependency model, is dynamic and based on the evaluation trace, similar to the manner in which Dynamic Program Slicing [9] extends classical Program Slicing. The value-based model records, in a detailed manner, which program parts actually contribute to the faulty behavior. By propagating value assumptions forward and backward via the standard diagnosis algorithms it provides better discrimination. In [5] an approach for combining program slicing and algorithmic debugging for debugging of procedural programs was introduced. The ideas of using test specifications and test results can be easily incorporated into our approach. In contrast to [5] we can change the debugging performance by changing the underlying model without changing the underlying algorithms.

The JADE debugger is a prototype system for research purposes and for demonstrating the underlying model-based techniques. Its main application interface can be seen in Figure 3. The system is not yet applicable in a real production environment. The incorporated Java evaluator that is used for computing variable values to be presented to the user during debugging implements only the basic Java functionality and some important classes and is far away from being JDK compliant. In addition, JADE assumes that the

Java source code for all involved classes is available. Both problems have to be tackled for a production version of JADE.

We have described the application of model-based technology in the building of a prototype intelligent debugger for Java programs. The approach is based on automatically building a formal internal model of the executed program, which a generic diagnosis engine then uses, together with observations of incorrect program output, for identifying possible sources of the error in the program. This combines modeling flexibility with the ability to reuse standard algorithms like measurement selection. Following an approach that we have followed in the domain of hardware design languages, the currently used model is purely dependency based for simplicity and quick diagnosis, and will be combined with a more detailed value-based model for improved discrimination capabilities.

References

[1] G. W. Bond and B. Pagurek. A Critical Analysis of "Model-Based Diagnosis Meets Error Diagnosis in Logic Programs". Tech. Rep. SCE-94-15, Carleton Univ., Dept. of Systems and Computer Engineering, Ottawa, Canada, 1994.

[2] Luca Console, Gerhard Friedrich, and Daniele Theseider Dupré. Model-based diagnosis meets error diagnosis in logic programs. In *Proc. IJCAI*, pages 1494–1499, Chambery, August 1993.

[3] J. de Kleer and B. C. Williams. Diagnosing multiple faults. *Artificial Intelligence*, 32(1):97–130, 1987.

[4] Gerhard Friedrich, Markus Stumptner, and Franz Wotawa. Model-based diagnosis of hardware designs. *Artificial Intelligence*, 111(2):3–39, July 1999.

[5] Peter Fritzson, Tibor Gyimothy, Mariam Kamkar, and Nahid Shahmehri. Generalized Algorithmic Debugging and Testing. In *Proc. Conf. on Programming Language Design and Implementation*, Canada, 1991.

[6] Peter Fröhlich and Wolfgang Nejdl. A Static Model-Based Engine for Model-Based Reasoning. In *Proc. 15th IJCAI*, Nagoya, Japan, August 1997.

[7] Daniel Jackson. Aspect: Detecting Bugs with Abstract Dependences. *ACM TOSEM*, 4(2):109–145, 1995.

[8] Bogdan Korel. PELAS–Program Error-Locating Assistant System. *IEEE TSE*, 14(9):1253–1260, 1988.

[9] Bogdan Korel and Jurgen Rilling. Applications of Dynamic Slicing in Program Debugging. In *Proceedings of the Third International Workshop on Automatic Debugging (AADEBUG-97)*, Linköping, Sweden, 1997.

[10] Cristinel Mateis, Markus Stumptner, and Franz Wotawa. Debugging of Java programs using a model-based approach. In *Proc. DX'99 Workshop*, Loch Awe, Scotland, 1999.

[11] Cristinel Mateis, Markus Stumptner, and Franz Wotawa. Locating bugs in Java programs – first results of the Java Diagnosis Experiments (Jade) project. In *Proceedings IEA/AIE*, New Orleans, 2000. Springer-Verlag.

[12] Antoine Missier, Spyros Xanthakis, and Louise Travé-Massuyés. Qualitative Algorithmics using Order of Growth Reasoning. In *Proc. ECAI 94*, pages 750–754, 1994.

[13] Mark Moriconi and Timothy Winkler. Approximate reasoning about the semantic effects of program changes. *IEEE Transactions on Software Engineering*, 16(9):980–992, September 1990.

[14] Raymond Reiter. A theory of diagnosis from first principles. *Artificial Intelligence*, 32(1):57–95, 1987.

[15] Markus Stumptner and Franz Wotawa. Debugging Functional Programs. In *Proc. 16th IJCAI*, Stockholm, Sweden, August 1999.

[16] Mark Weiser. Program slicing. *IEEE TSE*, 10(4):352–357, July 1984.