

Determination of the Topology of a Directed Network

Darin Goldstein
Computer Science Department
Cal State Long Beach
daring@cecs.csulb.edu

Abstract

We consider strongly-connected, *directed* networks of identical synchronous, finite-state processors with in- and out-degree uniformly bounded by a network constant. Via a straightforward extension of Ostrovsky and Wilkerson's Backwards Communication Algorithm[7], we exhibit a protocol which solves the Global Topology Determination Problem, the problem of having a root processor map the global topology of a network of unknown size and topology, with running time $O(ND)$ where N represents the number of processors and D represents the diameter of the network. A simple counting argument suffices to show that the Global Topology Determination Problem has time-complexity $\Omega(N \log N)$ which makes the protocol presented asymptotically time-optimal for many large networks.

1 Introduction

1.1 The Network Model. We consider strongly-connected *directed* networks of identical synchronous finite-state automata with in- and out-degree bounded by a constant. These automata are meant to model very small, very fast processors. The network itself has unknown topology and potentially unbounded size N . Throughout this paper, we use the term “-port” to refer to one of a number of unidirectional conduits through which constant-size messages may pass from one processor to another. An *in-port* to a processor will allow messages to flow unidirectionally in towards the processor. *Out-port* is defined similarly. We assume throughout that the number of in-ports and out-ports for each processor is uniformly bounded above by a network constant $\delta \geq 2$. The network is formed by connecting out-ports from processors to the in-ports of other processors with wires. Not all in-ports or out-ports of a given processor need necessarily be connected to other processors; however, any given processor must have at least one connected in-port and out-port. Note that even though the communication links are unidirectional, a pair of processors is allowed to be connected with two communication links, one in

either direction, simulating a bidirectional link.

We assume that each processor in the network is initially in a special “quiescent” state, in which, at each time-step, the processor sends a “blank” character through all of its out-ports. A processor remains in the quiescent state until a non-blank character is received by one of its in-ports.

The network has a global clock, the pulses between which each processor performs its computations. Processors synchronously, within a single global clock pulse, perform the following actions in order: read in the inputs from each of their in-ports, process their individual state changes, and prepare and broadcast their outputs.

The reason for modeling the processors by identical finite-state automata is simple. In practice, many network protocols are expected to run extremely fast. (One particular reason for this is that the network topology or size might change if the protocol takes too long thereby potentially rendering the computation obsolete. Obviously, if a processor is randomly added or removed from the topology of the network in the middle of the computation, a global topology determination is likely to produce an incorrect result.) Commonly, a memory access can take orders of magnitude longer than a simple state-change processor calculation. It is therefore assumed that the processors involved will not have time to access a large memory cache. The current technological trend is to merge the memory functions that one generally associates with a computing machine into the processor itself.

The protocol described below is presumed to begin when a certain processor is signaled by some outside source. We call this processor the *root*, and assume that every processor knows whether or not it is the root. A protocol ends when the root enters a special terminal state indicating that the computation has successfully completed. In our computational model, we calculate the time-complexity of a protocol in terms of the total number of global time steps between initiation and termination. Of course, the aim is to minimize this time-complexity.

DEFINITION 1.1. *Throughout this paper, we will use N to represent the total number of processors in the network. D will represent the diameter of the directed network.*

1.2 The Global Topology Determination Problem

1.2.1 Statement of the Problem. As previously mentioned, our computational model is designed to realistically simulate a large network of small and fast processors with only the capacity for reliable unidirectional communication. More specifically, we have a large network of powerful computers each equipped with a very fast communication processor (a communication device separate from the main processor of the computer which is presumably engaged in tasks other than simple communication protocols). These communication processors are modeled by the finite-state automata. (We think of finite-state machines as having a small/constant amount of memory which allows them to work faster than a larger, more complex machine.) One such computer initiates the protocol by nudging its communication processor out of quiescence thereby forcing the device to take on the role of the root. The protocol then commences via messages passing between the various communication processors. At each step of the protocol, the root is piping its computational transcript to the computer to which it is attached. By the time the protocol has completed execution (i.e. the root has re-entered quiescence after informing its master computer that the algorithm has completed), the root's computer has enough information in the form of input symbols read from its communication device to reconstruct the global topology of the network.

One point should be clarified: When considering explicit protocols on finite-state processors, an important characteristic of the processors is the ability to recognize whether their in-ports and out-ports are connected to other processors or not. The ability of a processor to recognize whether its in-ports are connected to other processors is called *in-port awareness*. A similar definition applies to *out-port awareness*. When dealing with practical applications it is natural to assume that processors have in-port awareness. (The quiescent (resting) state of a processor is to constantly send out the blank character b . If a processor is receiving \emptyset (nothing) and not b through one of its in-ports, the processor knows immediately that there cannot be a processor connected to that in-port.) However, out-port awareness is not necessarily a given in some communications systems. We will assume throughout this paper that processors have both in-port and out-port awareness as did Ostro-

vsky and Wilkerson in [7].

1.2.2 Previous Research. Mapping the global network topology is an extremely important primitive utilized for message routing as well as investigated for its intrinsic theoretical interest. The literature on the subject of message routing and topology mapping is immense. The most obvious example of practical network topology determination is Internet mapping. Internet mapping protocols are not in short supply. Mainwaring et al. in [5], for example, present and prove the theoretical correctness of a protocol for mapping the Myrinet system-area network at U.C. Berkeley. More ambitiously, R. Govindan and H. Tangmunarunkit [4] present and discuss a program called Mercator which performs "informed random address probing", a heuristic useful for mapping the entire Internet topology. And perhaps most impressively, in [1], Cheswick, Burch, and Branigan actually map the Internet topology almost entirely. (In fact, they even perform an analysis of the NATO bombing campaign on Yugoslavia's network connectivity.) Numerous other mapping protocols and results can be found in these respective bibliographies.

Of course, the Internet is commonly considered a bidirectional network with a fairly predictable overall topology (from a theoretician's point of view). We will focus on *directed* networks of totally unknown topology and potentially unbounded size, and the processors are assumed to be finite-state. The fact that our network assumptions are so general make our results applicable to other, more specific networks (e.g. any of the above references). However, our solution might not be the most efficient from a practical point of view for these other specific network types; this is especially true if one has the foreknowledge that the network in question has particular properties such as bidirectionality, greater processor memory, or an easy-to-map hypercube topology. On the other hand, networks of the more general kind occur more often than one might think (e.g. GPS satellites, encrypted one-way radio military networks, bidirectional networks with in-port or out-port shutdown failures at individual processors, etc.). Network restrictions such as unidirectional communication and finite-state processors make outlining a general topology mapping protocol a nontrivial exercise. Luckily, Even, Litman, and Winkler's "snake" data structure [2] (used in Ostrovsky and Wilkerson's BCA in [7] and the equivalence proof in [3] and modified for use in the protocol below) is virtually tailor-made for the purpose.

2 Data Structures

2.1 Speed. The protocols about to be presented make use of several computational constructs, each of

which is assigned a certain characteristic “speed.” (The “speed” concept can be referenced as far back as [6].) This is *not* to say that certain messages move faster through the network than others. All computations and outputs are strictly synchronous with respect to the global network clock.

In the protocol that follows, the speeds that we utilize are speed-1 and speed-3. The method by which we implement a speed is as follows: A speed-1 construct will enter a processor. It will then remain there for 3 global clock ticks. At the third clock tick, it will proceed along its designated path. Similarly, a speed-3 construct will wait only 1 global clock tick. Thus, in reality, the implementation specifies that a speed-1 construct moves 3 times slower than a speed-3 construct.

2.2 Tokens. Tokens are the simplest data structure possible on networks of finite-state machine. They should be thought of as markers that can be passed from one processor to another via the edges of the network. The token concept has been in use since the first solution to the Firing Squad Synchronization Problem for the bidirectional line [6]. The definitions we give below of “breadth-first” and “loop” tokens were first outlined in [3] though the behaviors were also utilized in [2, 7].

We employ two main varieties of tokens. *Breadth-first tokens* can be thought of as moving within a “breadth-first-search tree” in the following sense: We arrange it so that each relevant processor in the network has a “parent” marker associated with one of its in-ports. (The method by which various breadth-first-search trees are constructed by snakes, as well as how each processor designates its “parent” in-port, is discussed in Section 2.3.2.) We then declare that a breadth-first token will only be accepted by a given processor when either (a) the processor creates the token, in which case the processor will not have a parent in-port, or (b) the token comes through the processor’s parent in-port. If a breadth-first token is received through a non-parent in-port or by a quiescent processor, it is ignored. Breadth-first tokens are passed out of every out-port; thus, breadth-first tokens multiply in number as time goes on (as long as they stay within the confines of the breadth-first-search tree.) In summary, if a breadth-first token is created at the root of its associated breadth-first tree, then t time steps later there will be a token at each processor that is a distance of t from the root, and none elsewhere. (If the tree has length less than t , of course, there will be no tokens anywhere.)

Loop tokens travel along a specified marked loop within the network. (How loops get marked is described in Section 2.4.) A processor on the loop that receives a loop token simply passes it to the next processor on

the loop. Thus, t time steps after its creation, a loop token will be t processors away from the processor that created it, along the marked loop. When any loop token reaches its creator processor, it is absorbed (i.e., not sent around again).

Note that tokens can only carry along with them a constant (very small) amount of information since they are only of constant size. The next data structure takes care of this problem.

2.3 Snakes. Our description of the snake data structure closely follows that in [3]. The concept of a data-carrying *snake* was invented by Even, Litman, and Winkler in [2]. Snakes are the solution to the problem of the limited data-carrying capabilities of tokens. A snake is capable of carrying an arbitrarily large amount of data, but for this reason, it must reside in a collection of adjacent processors rather than a single processor.

A “snake” is a string – which may be arbitrarily long – made up of an alphabet of $2(\delta^2 + \delta) + 1$ characters, namely $\delta^2 + \delta$ head characters, $\delta^2 + \delta$ body characters, and a unique tail character. (Recall that δ is a fixed constant of the network.) The characters comprising the string are stored in adjacent processors, one character per processor. These characters encode a path by specifying a series of in- and out-ports. (Note that a token could never do such a thing, since a path in the network can grow arbitrarily long.)

We require two main snake types, which we call growing and dying. *Growing snakes* are used to generate encoded paths of the network, and *dying snakes* are used to mark encoded paths. Our protocol requires two kinds of each of the two snake types; specifically, we will need out-growing, in-growing, out-dying, and in-dying snakes. “Out” and “in” are meant as a mnemonic; out-snakes are generated at the root and proceed outward from it, while in-snakes are generated elsewhere and trigger some action when they reach the root. Out-growing, in-growing, out-dying, and in-dying snakes will be referred to as OG-snakes, IG-snakes, OD-snakes, and ID-snakes in what follows.

Each of the four kinds of snake gets its own alphabet of characters to describe it; this allows processors to determine with which kind of snake they are dealing. We will spend a section on each type, elucidating its respective behavior. First, we need to go over some general rules common to all snake types.

2.3.1 General Snake-handling Rules.

- All snakes are speed-1.
- Snakes of different types do not interact. A processor can handle different snake types simulta-

neously without getting confused because snake types are distinguished by their alphabets. Note that this does not impose arduous memory constraints upon the processors (which are finite-state machines) since there is only a constant number of snake types.

2.3.2 Growing Snakes. Growing snakes function as information generators. We define the *initiator* to be the processor from which the growing snakes first emanate. The *terminator* is defined to be the processor that the snakes are attempting to reach. Growing snakes grow in a breadth-first manner; the first growing snake to reach the terminator processor will have encoded within its body a minimal-length path from the initiator to the terminator. Upon reaching the terminator, a growing snake head might then initiate some further action based on the protocol and the state of the terminator. The rules for handling growing snakes are outlined below; the rules for handling in-growing and out-growing snakes are identical (just replace “IG” with “OG”). We assume that we are using in-growing snakes in most of the discussion below for concreteness.

- First, the head characters of the baby growing snakes are generated by the initiator. This processor sends an IG-snake head character out of every out-port during the first time step. The particular head character to be sent will correspond to the out-port from which it is being sent. For every i between 1 and δ , the growing head snake character $IGH(i, *)$ will be sent through out-port i . When a processor receives any growing snake character with $*$ as its second parameter (and this applies to body as well as head characters), the processor notes the in-port j through which the character arrived and changes the $*$ to j . Thus, when the $IGH(i, *)$ is received, it is changed to $IGH(i, j)$ where j is the number of the receiving in-port. During the next time step, the initiator will send a tail character IGT through every out-port. Thus a baby snake is born.
- When a processor receives an in-growing snake character (again, for concreteness) for the first time, it marks itself *IG-visited*, and marks the in-port through which the growing snake character was passed as its *IG-parent*¹. (These marks will be used later by certain breadth-first tokens; see Section 4.2.1.) Only this first IG-snake will be

allowed to pass through the processor; all other IG-snake characters will be ignored. Thus, an IG-snake will carve out a breadth-first-search tree. Growing snake characters are periodically removed from the network. Until this removal occurs, however, each growing snake carves out a breadth-first-search tree.

- When a processor receives a non-tail IG-snake character, it simply sends this character through all out-ports during the next time step. Once the processor sends the character out, it need not retain it in “memory.” (In this way, the processor simply passes the head and body through every out-port. Thus arbitrarily long paths can be generated.)
- Once a processor receives the tail of an IG-snake, instead of simply passing it through like the other body characters, the processor generates a new body character. For all i between 1 and δ , it simultaneously sends the character $IG(i, *)$ through out-port i ; thus a new body character is generated to mark the current processor’s position in the path. Only after this new character is passed along does the processor send the tail through. Note that the $*$ is changed to reflect the appropriate in-port when the body character is received by the next processor in turn.

2.3.3 Dying Snakes. Dying snakes function as path markers. After a path is generated by the growing snakes, it is the responsibility of the dying snake to mark the generated path so that the processors on it will know (a) that they lie along a special path and (b) which in-port and out-port they should use for funneling information along the path. In our protocol, ID-snakes will be formed by converting the characters of an OG-snake into ID-snake characters as they pass through one particular processor; OD-snakes will be created from ID-snakes in a manner to be described in Section 4.2.1. The rules for handling ID-snakes are outlined below; the rules for handling OD-snakes are identical (just replace “ID” with “OD”), except where noted.

- An ID-snake will mark a path generated by an OG-snake (see Section 2.3.2); thus, since an OG-snake carves out a breadth-first-search tree, the path will never self-intersect. Similarly, neither will a path that an OD-snake is to mark. However, the concatenation of the two paths (which, in our protocol, will always be a loop that includes the root) may self-intersect; any processor will appear at most twice on the concatenation. We will, eventually, want to consider the concatenation as

¹If two or more IG-snakes arrive simultaneously, the one arriving through the lowest-numbered in-port is deemed “first.”

a whole; to this end, we imbue each processor with two “predecessor in-ports” (numbered 1 and 2) and two “successor out-ports” (ditto).

- Whenever a processor receives the head character $IDH(i, j)$ of an ID-snake, it sets predecessor in-port #1 equal to the number of the in-port through which it received the character, and sets successor out-port #1 equal to i . These two values indicate the two edges of the path incident to the processor. OD-snakes work identically, except that they use predecessor in-port #2 and successor out-port #2. The head character is then discarded (not sent through any out-port).
- If the next ID-snake character that the processor receives through the predecessor in-port is $ID(i', j')$, it gets sent through the successor out-port as $IDH(i', j')$. (In other words, the next ID-snake body character that comes through gets converted to the new head.) The processor then passes all further ID-snake characters received through its predecessor in-port to its successor out-port exactly as received. If the next character happens to be a tail, then it gets sent through the successor out-port as is. In our protocol, ID-snakes will be converted into OD-snakes at the root; this will provide an exception to these rules, for at the root all ID-snake characters are converted into OD-snake characters instead. In addition, as might be expected, the root will set predecessor in-port #1 and successor out-port #2 appropriately as the dying snakes go through; its other two ports will not be needed, as we will show in Section 4.2.1.
- An ID-snake only propagates along the path it is marking, and a maximum of one will be in the network at any given time, so we need not worry about ID-visited markings.

2.4 Marked loops. As mentioned in Section 2.3.3, we will be using dying snakes to mark certain loops (not necessarily simple) that include the root. We will refer to this structure repeatedly throughout the paper, and thus make the following definition:

DEFINITION 2.1. *A marked loop will be a loop marked by dying snakes in the manner described in Section 2.3.3. The root must be one of the processors on the loop. The loop may or may not be simple, but no processor or edge can appear more than twice on it.*

Each processor will have its predecessor and successor port (or, if necessary, ports) set by the dying snakes. A processor with only predecessor in-port #1

set will only accept a loop token through that in-port; it will subsequently pass the token through successor out-port #1². Similarly, a processor with only predecessor in-port #2 set will only accept a loop token through that in-port; it will subsequently pass the token through successor out-port #2. Finally, a processor with both predecessor in-ports set will initially accept a given loop token only through predecessor in-port #1 (it will pass the token through successor out-port #1, of course); it then waits for the token to come through predecessor in-port #2 (at which point it passes the token through successor out-port #2); it then will expect the next such loop token through predecessor in-port #1 again.

We will hereon refer to the predecessor in-port (resp. corresponding successor out-port) through which a loop processor awaits a loop token as the *appropriate* predecessor in-port (resp. successor out-port).

3 The Global Topology Determination Algorithm

3.1 Description of the Algorithm. In the discussion to follow, we will assume that two auxiliary protocols, the Backwards Communication Algorithm (BCA) and the Root Communication Algorithm (RCA), are available for use. The BCA is a method for sending information “backwards” along a unidirectional edge in the network, and the RCA is a method for communication information from any given node to the root. We defer more complete descriptions of each until Sections 4.1 (BCA) and 4.2 (RCA).

After initiation by its master computer, the root releases a DFS (Depth First Search) token through its lowest-numbered connected out-port. This token performs a depth first search of the entire network remembering along the way through which out-port it has been most recently passed and through which in-port it was most recently received. (The DFS token is to be thought of as having the same basic structure as a snake character with two entries where in-port and out-port labels can be stored.) The information stored in this token is conveyed to the root as the depth-first-search progresses.

We assume that the reader is somewhat familiar with the mechanics of depth-first search on directed graphs. We will give a brief overview of the depth-first search using finite-state processors. In the following discussion, a “forward edge” refers to an existing edge of the network representing a path along which messages are passed unidirectionally. The reason we even bother

²Once again, the root will provide an exception to this rule; it will accept a loop token only through predecessor in-port #1, but will pass it through successor out-port #2;

making this distinction is that because of the BCA, we have a method of passing information *backwards* through an edge. When a piece of information gets passed through a legitimate edge of the network, we say that it gets passed *forward* through the edge.

To perform the depth-first search, any given processor, after receiving the DFS token for the first time³, notes the in-port through which it received the DFS token. The processor also marks that in-port as its parent and then passes the DFS token out its lowest-numbered connected out-port. After the processor gets the DFS token back *via the BCA*⁴, it marks that out-port finished and sends the DFS token out of the lowest-numbered unfinished connected out-port, and so on. When all of the processor's out-ports are finally finished, the processor sends the DFS token back through its parent in-port via the BCA. Once the root has finished all of its out-ports, the depth-first search is over.

The root is updated as to the progress of the protocol via the following. Upon receipt of the DFS token, a processor initiates one of the following two tasks. Once the task is completed, the DFS token is passed on according to the rules of depth-first-search outlined above. (As indicated above, whenever the DFS token needs to move backwards along an edge of the network, it uses the BCA.)

- If the token was *not* received through use of the BCA, the processor performs the RCA using the FORWARD token. We assume that there are δ^2 possible FORWARD tokens. The FORWARD token that gets sent depends on which out-port sent the DFS token and which in-port received the DFS token. For example, if the DFS token was passed out of out-port 4 of one processor and into in-port 1 of another, then FORWARD token (4, 1) is sent. (The FORWARD token has the same basic structure as a snake character.)
- If the DFS token was received via a backwards edge (i.e. if the token was passed to the processor by the BCA), the processor performs the RCA using the BACK token.

The algorithm terminates when the root has completed the depth first search of the network (i.e. finished all of its out-ports).

³If a processor is receiving the DFS token for the first time, it must be through a forward edge of the network.

⁴A processor may get the DFS token back through a forward edge of the network after it already has marked its parent in-port. In that case, the processor would use the BCA to send the DFS token back since a processor never wants more than one parent.

What is the master computer's strategy for mapping the network given the computational transcript of its communication processor at the root? We will describe the strategy as if the computer were drawing a topological map as the algorithm was proceeding. The computer always keeps track of the processors in the network that have performed previous RCA's, allocating them names as new processors are "discovered" by the algorithm. (Recall that the computer at the root has the ability to assign processors unique names even though the communicating devices at the processors themselves cannot.) It will keep a stack of processor positions as well. When a processor performs an RCA with a FORWARD token, the computer pushes it onto the stack. The processor at the top of the stack then points to the processor that has most recently performed a RCA. (If the root has just initiated the Global Topology Determination Protocol then we consider the root itself as having performed the last RCA; the stack will initially consist of only the root.) Whenever an RCA is run, the computer notes the characters of the IG-snake that passes through the root as it is converted to an OG-snake (see Lemma 4.1). From the characters of the IG-snake, the root computer can accurately map both the in-ports and the out-ports of the canonical shortest path to the current processor A , the processor running the RCA. Because the protocol is deterministic and always produces the same canonical shortest path from any given processor A to the root and back again, the computer can tell whether the current processor A has already been marked on the map. If it has not yet been marked on the map, the computer marks it and creates a name for it. At the end of the RCA, the computer notes whether a FORWARD or BACK token is being passed around the loop. If it is a FORWARD token, then the computer should draw a directed arrow from the top processor on the stack to the current processor A through the appropriate out-port and in-port. Afterwards, the computer pushes processor A onto the stack. (Recall that a FORWARD token indicates that the depth first search has moved forward along an edge.) If it is a BACK token, the computer simply pops the top processor off the stack. Note that the top processor on the stack tracks the position of the DFS token at any given point in time.

4 Auxiliary Protocols and Correctness Proofs

4.1 The Backwards Communication Algorithm. The Backwards Communication Algorithm (BCA) first appeared in [7] and accomplishes the following: Assume there is a directed edge from processor A to processor B in the network. The BCA is a method for sending a message from processor B to processor A

(backwards through the directed edge) such that A gets the message, B knows when A has gotten the message, and at the end of the transaction, the rest of the graph is left undisturbed⁵. The running time for each use of the BCA is $O(D)$.

DEFINITION 4.1. *Given two distinct processors in the network, processor A and processor B , we define the canonical shortest path from processor A to processor B to be the unique path along which the first growing snake released from processor A that survives to reach processor B would travel.*

4.2 The Root Communication Algorithm. In this section we will outline another algorithm, which we call the Root Communication Algorithm (RCA), based on the idea of Ostrovsky and Wilkerson’s BCA, which we will use as part of the Global Topology Determination Protocol presented in Section 3. The RCA accomplishes the following: processor A communicates a message to the root such that the root gets the message, processor A is aware of the completion of the algorithm, the computer at the root is able to reconstruct the sequence of in-ports and out-ports along the canonical shortest paths leading from the the root to processor A and from processor A to the root, and at the end of the transaction, the graph is left undisturbed. This auxiliary algorithm will be used to send one of the two signals FORWARD or BACK to the root and to allow the root computer to track the movement of the DFS token. The Global Topology Determination Protocol is guaranteed to only be running the RCA at a single processor at any given time. Throughout this section, we will assume that processor A is the processor that wishes to communicate with the root.

4.2.1 The Steps of the Root Communication Algorithm. For the sake of brevity, we will abbreviate the snake types.

1. Processor A becomes aware that it wishes to communicate with the root (via a process outlined in the steps of the Global Topology Determination Algorithm in Section 3) and sends IG-snakes to search for the root. Any processor receiving an IG-snake character for the first time marks itself as “IG-visited”, thus preventing any subsequent IG-snakes from entering it. It will also designate the

in-port from which it received the IG-snake as its “IG-parent” in-port. These markings will not be cleared until the release of KILL tokens (in step 4); hence, the IG-snakes carve out an IG-breadth-first-search tree.

2. Upon receipt of the head of the first IG-snake to reach it, the root performs two actions simultaneously. First, the root closes itself off to all other IG-snakes, ignoring any that attempt to enter. The root will accept no further IG-snakes during this execution of the algorithm. Second, the root begins to convert the IG-snake to an OG-snake which it broadcasts out all out-ports. (To convert an IG-snake to an OG-snake, the root simply converts the IG-snake characters it receives as input to OG-snake characters when they are sent out.) When the root receives the tail of the IG-snake that it is converting, it simply holds onto the tail and continues growing the OG-snake normally: The root holds onto the tail character while it sends out the character $OG(i, *)$ out of each of its out-ports for every i between 1 and δ . Only after this character is broadcast does the root send out the tail of the snake as OGT . OG-snakes leave “OG-visited” and “OG-parent” markings (and thus create an OG-breadth-first-search tree) similar to the IG-snakes discussed in Step 1. However, the OG-snakes do not respect the IG-breadth-first-search tree and are therefore guaranteed to make it back to processor A .
3. When processor A receives the first OG-snake head that survives, processor A closes itself off to any subsequent OG-snakes and converts the OG-snake to an ID-snake (note that processor A must therefore eat the head character of the OG-snake as if it were an ID-snake character, then send the rest of the snake through the appropriate out-port). The ID-snake then marks the path from processor A to the root. Eventually the root receives the head of an ID-snake and converts it to an OD-snake as previously described. This OD-snake then marks the path from the root back to processor A which will only receive the tail character ODT . At this point in the protocol, there is a marked path in the network from processor A to the root and back again.
4. As soon as processor A receives the tail of the OD-snake, processor A performs two tasks simultaneously. First, it releases a speed-3 breadth-first KILL token. The function of the KILL token is to completely eradicate all traces of growing snake characters in the network; both IG- and

⁵It is important to note that the names A and B are just names for the reader’s convenience. Processors cannot all simultaneously assign themselves unique names because of their finite-stateness. Processor B only recognizes Processor A as the processor that is on the other end of one of Processor B ’s in-ports.

OG-snake characters and markings are erased upon contact with a KILL token. KILL tokens are ignored by those processors that do not have any growing snake markings or characters in their memory. (KILL tokens do not affect the marked path.) Second, processor A releases a speed-1 loop token. This token will either be a FORWARD or BACK token depending on the current state of the network. Upon reception of the FORWARD/BACK token, processor A is guaranteed that one time step later, there will be no further growing snake characters or KILL tokens percolating uselessly through the network.

5. Processor A finally releases a speed-3 UNMARK token around the marked path. Each processor the old marked loop, upon receiving the token through its appropriate predecessor in-port, passes the token through the appropriate successor out-port, then forgets those predecessor and successor designations. Upon reception of this UNMARK token, the root reopens itself to IG-snakes. After the token makes it all the way around the marked path back to processor A , processor A reopens itself to OG-snakes and terminates the algorithm.

4.2.2 Proof of Correctness of the RCA. In this section, we prove the non-trivial claims made about the RCA presented in Section 4.2.1.

LEMMA 4.1. *The root's master computer is able to determine the canonical shortest paths leading to and from processor A by the completion of the algorithm.*

Proof. Note that the canonical shortest path from processor A to the root is unique and is encoded in the body of the first (and only) IG-snake to safely reach the root. Thus, to track this path, the master computer can simply read off the in-ports and out-ports encoded in the body of the IG-snake as it is converted to an OG-snake in Step 2. Similarly the canonical shortest path from the root to processor A is also unique and is encoded in the body of the ID-snake that reaches the root in Step 3. The master computer can again simply read off the relevant in-ports and out-ports as the ID-snake is converted to an OD-snake.

LEMMA 4.2. *After processor A terminates the algorithm in Step 5, the network is left completely undisturbed by any data construct created by the algorithm (snake characters/markings, tokens, etc.).*

Proof. Because KILL tokens travel three times faster than snakes, it is obvious that the tokens will eventually catch up with and eliminate the growing snakes. To

see that the KILL tokens will catch up precisely when claimed (see step 4), note that if L is the length of the current marked loop in the network, then the snake heads have at most a $2L$ head start. This implies that the KILL tokens will catch up with the snake heads after a speed-1 token (i.e. the FORWARD/BACK token) makes it around the loop once.

LEMMA 4.3. *Each execution of the RCA by any given processor A takes time $O(D)$.*

Proof. By inspection of the steps, the running-time is proportional to the length of loop marked by the algorithm: $d(A, \text{root}) + d(\text{root}, A)$. This quantity is itself trivially $O(D)$.

4.3 Correctness Proof of the Global Topology Determination Algorithm

LEMMA 4.4. *The Global Topology Determination Algorithm terminates in time $O(ND)$.*

Proof. Each processor in the network performs at most δ RCA's and at most δ BCA's. The running time for each of these subalgorithms is $O(D)$ and the result follows.

THEOREM 4.1. *The computer at the root of a network performing the Global Topology Determination Algorithm accurately maps the given directed network.*

Proof. First, we claim that any time a FORWARD token is noted by the root computer, the two processors between which it draws the directed arrow have both already been mapped. Every time an RCA is run, processor A must be mapped before the token is sent out because a IG-snake must be sent through the root before a FORWARD token. The previous processor (i.e. the processor on top of the stack) has already been mapped by a previous execution of the RCA and thus a FORWARD token can always be traced between two processors already on the network map at the time it is sent. Now we note that the DFS token must be sent forward through every edge of the network and hence a FORWARD token is sent for every edge of the network. Thus all edges get accurately mapped.

5 The Lower Time Bound

LEMMA 5.1. *Let $G(N)$ be the number of bounded-degree strongly connected networks of N processors and diameter less than or equal to $2\log N + 1$ with distinct topologies. Then there exists some constant C such that, for large enough N , $G(N) \geq N^{CN}$. (i.e. There are a great many networks with small diameter.)*

Proof. We only present a quick justification and leave the interested reader to fill in the details. Consider the family of networks that consist of a full binary tree emanating from a single node with bidirectional edges (i.e. unidirectional edges in both directions) with a simple loop that includes every processor on the bottom level of the tree. Note that all such networks are of bounded-degree and strongly connected. Every rearrangement of the processors included in the loop on the bottom levels yields a distinct topology. A simple counting argument suffices to complete the proof.

We make the convention that the processors' input/output set is called I . The number of elements of the set I is $|I|$.

LEMMA 5.2. *For any given algorithm, after x global clock ticks, the root can have had one of a maximum of $|I|^{\delta x}$ possible computational transcripts.*

THEOREM 5.1. *Any algorithm which solves the Global Topology Determination Problem has a time-complexity lower bound of $\Omega(N \log N)$.*

Proof. Fix an algorithm which solves the Global Topology Determination Problem. Let us assume that the algorithm terminates on graphs with N processors in less than or equal to $T(N)$ global clock ticks.

In order for the root to distinguish between different global topologies, for any given network size N , there must be at least as many computational transcripts as there are distinct network topologies. Otherwise, by the pigeonhole principle, two distinct network topologies would have to be distinguished by exactly the same computational transcript which is, of course, impossible.

By Lemma 5.1, we know that for large enough N , there exists a constant C such that the number of distinct network topologies is greater than or equal to N^{CN} . By Lemma 5.2, the number of possible computational transcripts the root can have had is at most $|I|^{\delta T(N)}$. Thus we get:

$$|I|^{\delta T(N)} \geq N^{CN} \Rightarrow T(N) = \Omega(N \log N)$$

References

- [1] Hal Burch and Bill Cheswick. Mapping the internet. *IEEE Computer*, 32(4):97–98, 102, 1999.
- [2] S. Even, A. Litman, and P. Winkler. Computing with snakes in directed networks of automata. *J. Algorithms*, 24(1):158–170, 1997.
- [3] D. Goldstein and N. Meyer. The wake up and report problem is asymptotically time-equivalent to the firing squad synchronization problem. In *Proceedings of the Thirteenth Annual ACM-SIAM Symposium on Discrete Algorithms*, pages 578–587, San Francisco, CA, 6–8 January 2002.
- [4] Ramesh Govindan and Hongsuda Tangmunarunkit. Heuristics for internet map discovery. In *INFOCOM (3)*, pages 1371–1380, 2000.
- [5] A. Mainwaring, B. Chun, S. Schleimer, and D. Wilkerson. System area network mapping. *ACM Symposium on Parallel Algorithms and Architectures*, 1997.
- [6] M. Minsky. *Computation: Finite and Infinite Machines*. Prentice-Hall, Englewood Cliffs, NJ, 1967.
- [7] R. Ostrovsky and D. Wilkerson. Faster computation on directed networks of automata. In *Proceedings of the Fourteenth Annual ACM Symposium on Principles of Distributed Computing*, pages 38–46, Ottawa, Ontario, Canada, 2–23 August 1995.