

INSTITUT NATIONAL DE RECHERCHE EN INFORMATIQUE ET EN AUTOMATIQUE

Implémentation asynchrone de détecteurs de fautes sans connaître les participants et en présence d'une connectivité partielle

Pierre Sens — Luciana Arantes — Mathieu Bouillaguet — Véronique Martin — Fabíola Greve

 N° 6088 — version 2

version initiale Janvier 2007 — version révisée Décembre 2013

Thème COM

l apport de recherche



Implémentation asynchrone de détecteurs de fautes sans connaître les participants et en présence d'une connectivité partielle

Pierre Sens , Luciana Arantes , Mathieu Bouillaguet , Véronique Martin * , Fabíola Greve †

Thème COM — Systèmes communicants Projets Regal

Rapport de recherche n° 6088 — version 2 — version initiale Janvier 2007 — version révisée Dcembre 2013 — 23 pages

Résumé: Cet article aborde le problème de la détection de fautes dans les réseaux dynamiques de type MANET. Les détecteurs de fautes non fiables fournissent des informations sur les processus défaillants. Ils permettent de résoudre le consensus dans les réseaux asynchrones. Cependant, la plupart des détecteurs considère un ensemble connu de processus interconnectés par un réseau complètement maillé. Des telles hypothèses ne sont pas réalistes dans les environnements dynamiques. Généralement, les implémentations des détecteurs reposent sur des temporisateurs dont les bornes sont particulièrement difficiles à déterminer dans le contexte des réseaux dynamiques. Cet article présente une implémentation asynchrone de détecteurs de défaillances adaptée aux environnements dynamiques. Nous prouvons que notre algorithme permet d'implémenter un détection de classe $\Diamond S$ lorsque des propriétés sur la vitesse relative des transmissions et la connectivité sont satisfaites par le réseau sous-jacent.

Mots-clés: dtecteurs de fautes, algorithmes rpartis, rseaux dynamiques

 $^{^{\}ast}$ LIP6 - University of Paris 6 - INRIA

[†] DCC - Computer Science Department / Federal University of Bahia

Asynchronous Implementation of Failure Detectors with partial connectivity and unknown participants

Abstract: We consider the problem of failure detection in dynamic networks such as MANETs. Unreliable failure detectors are classical mechanisms which provide information about process failures. However, most of current implementations consider that the network is fully connected and that the initial number of nodes of the system is known. This assumption is not applicable to dynamic environments. Furthermore, such implementations are usually timer-based while in dynamic networks there is no upper bound for communication delays since nodes can move. This paper presents an asynchronous implementation of a failure detector for unknown and mobile networks. Our approach does not rely on timers and neither the composition nor the number of nodes in the system are known. We prove that our algorithm can implement failure detectors of class $\Diamond S$ when behavioral properties and connectivity conditions are satisfied by the underlying system.

Key-words: failure detectors, distributed algorithms, dynamic networks

1 Introduction

The distributed computing scenario is rapidly evolving for integrating unstructured, self-organizing and dynamic systems, such as peer-to-peer, wireless sensor and mobile ad-hoc networks. Nonetheless, the issue of designing reliable services which can cope with the high dynamism of these systems is a challenge.

Failure detector is a fundamental service, able to help in the development of fault-tolerant distributed systems. Its importance has been revealed by Chandra and Toueg who proposed the abstraction of unreliable failure detectors in order to circumvent the impossibility result of the consensus problem in an asynchronous environment [FLP85, CT96]. Unreliable failure detectors, namely FD, can informally be seen as a per process oracle, which periodically provides a list of processes suspected of having crashed. In this paper, we are interested in the class of FD denoted $\Diamond S$. Chandra and Toueg proved that by adding FD of class $\Diamond S$ to an asynchronous system, it is possible to deterministically solve the consensus problem (with the additional assumption that a majority of processes are correct).

This paper focuses on FD for mobile and unknown networks, such as mobile ad-hoc networks (MANETs). This kind of network presents the following properties: (1) a node does not necessarily know all the nodes of the network. It can only send messages to its neighbors, i.e., those nodes that are within its transmission range ¹; (2) message transmission delay between nodes is highly unpredictable; (3) the network is not fully connected which means that a message sent by a node might be routed through a set of intermediate nodes until reaching the destination node; (4) a node can move around and change its transmission range.

Most of current implementations of failure detectors are based on an all-to-all communication approach where each process periodically sends a heartbeat message to all processes [LFA00, SM01, DT00]. As they usually consider a fully connected set of known nodes, these implementations are not adequate for dynamic environments for the reasons explained above. Furthermore, they are usually timer-based, assuming that eventually some bound of the transmission will permanently hold. Such an assumption is not suitable for dynamic environments where communication delays between two nodes can vary due to mobility of nodes. In [MMR03], Mostefaoui et al. have proposed an asynchronous implementation of FDs which is timer-free. It is based on an exchange of messages which just uses the value of f (the maximum number of processes that can crash) and n (the number of nodes in the system). However, their computation model consists of a set of fully connected initially known nodes. Some recent works have been proposed which deals with the scalable nature of dynamic systems [LFA00, GCG01, BMS03]. Nonetheless, few of them tolerate mobility of nodes [FT05, TTS04] and they are all timer-based.

This paper presents a new asynchronous FD algorithm for dynamic systems of mobile and unknown networks. It does not rely on timers to detect failures and no knowledge about the system composition nor its cardinality is required. Yet, it has some interesting features

¹The concept of range models, for instance, homogeneous radio communication in MANETs.

that allow for scalability. The detection of process failures is based only on a local perception that the node has on the network and not on global exchanged information.

The basic principle of our FD is the flooding of failure suspicion information over the network. Initially, each node only knows itself. Then, it periodically exchanges a QUERY-RESPONSE pair of messages with its neighbors, that is, those nodes from which it has received a message previously. Then, based only on the reception of these messages and the partial knowledge about the system membership (i.e., its neighborhood), a node is able to suspect other processes or revoke a suspicion in the system. This information about suspicions and mistakes is piggybacked in the QUERY messages. Thus, as soon as the underlying system satisfies an f-covering property, suspicions and mistakes are propagated to the whole network. The f-covering property ensures that there is always a path between any two nodes of the network, in spite of f faults (f < n).

Moreover, if the processes in the system satisfy some behavioral properties, our algorithm implements the failure detectors properties of the class $\Diamond S$. Four behavioral properties have been defined. The *membership property* states that, in order to be known in the system, a node should interact (by sending messages) at least once with some others. The *mobility property* states that a moving node should reconnect to the network longtime enough in order to update its state regarding failure suspicions and mistakes. The *responsiveness property* states that after a given time, communication between some node in the system and its neighborhood is always faster than the other communications of this neighborhood. Finally, the *mobility responsiveness property* states that at least one correct node in the system does satisfy the responsiveness property and that its neighborhood is composed of non-moving nodes.

The rest of the paper is organized as follows. Section 2 presents Chandra-Toueg's failure detectors. Section 3 defines the computation model. In Section 4, our asynchronous failure detector algorithm is presented considering that nodes do not move. Section 5 describes how the algorithm can be extended to support mobility of nodes. Simulation performance results are shown in Section 6, while some related work are briefly described in Section 7. Finally, Section 8 concludes the paper.

2 Chandra-Toueg's Failure Detectors

Unreliable failure detectors provide information about the aliveness of processes in the system [CT96]. Each process has access to a local failure detector which outputs a list of processes that it currently suspects of having crashed. The failure detector is *unreliable* in the sense that it may erroneously add to its list a process which is actually correct. But if the detector later believes that suspecting this process is a mistake, it then removes the process from its list. Therefore, a detector may repeatedly add and remove the same process from its list of suspected processes.

Failure detectors are formally characterized by two properties. *Completeness* characterizes its capability of suspecting every faulty process permanently. *Accuracy* characterizes its capability of not suspecting correct processes. Our work is focused on the class of *Eventually*

Strong detectors, also known as $\lozenge S$. This class contains all the failure detectors that satisfy (1) Strong completeness: there is a time after which every process that crashes is permanently suspected by every correct process; (2) Eventual weak accuracy: there is a time after which some correct processes are not suspected by any correct process.

3 Model

We consider a dynamic distributed system consisting of a finite set Π of n > 1 mobile nodes, namely, $\Pi = \{p_1, \dots, p_n\}$. Contrarily to a static environment, in a dynamic system of mobile unknown networks, processes are not aware about Π and its cardinality n. Thus, they know only a subset of processes in Π . There is one process per node and they communicate by sending and receiving messages via a packet radio network. There are no assumptions on the relative speed of processes or on message transfer delays, thus the system is asynchronous. A process can fail by crashing. A correct process is a process that does not crash during a run; otherwise, it is faulty. Let f denote the maximum number of processes that may crash in the system (f < n). We assume that f is known to every process. To simplify the presentation, we take the range \mathcal{T} of the clock's tick to be the set of natural numbers. Processes do not have access to \mathcal{T} : it is introduced for the convenience of the presentation.

The system can be represented by a communication graph G(V, E) in which $V \subseteq \Pi$ represents the set of nodes and E represents the set of logical links. Nodes p_i and p_j are connected by a link $(p_i, p_j) \in E$ iff they are within their wireless transmission range. In this case, p_i and p_j are considered 1-hop neighbors, belonging to the same neighborhood. The topology of G is dynamic. Links are considered to be reliable: they do not create, alter or lose messages. Then, a message m broadcast by p_i is heard by all correct processes in p_i 's neighborhood. Communications between 1-hop neighbors are either broadcast or point-to-point.

When a node moves, we consider that it is separated from G. Afterwards, when it stops moving and reconnects to the network, it is reinserted to G. A node can keep continuously moving and reconnecting, or eventually it crashes. Nonetheless, a correct moving node will always reconnect to the network. A moving node is one that is separated from G and a non-moving node is connected to G. Let p_m be a moving node. We consider that p_m is not aware about its mobility. Thus, it cannot notify its neighbors about its moving. In this case, for the viewpoint of a neighbor, it is not possible to distinguish between a moving or a crash of p_m . During the moving, p_m keeps its state, that is, the values of its variables.

Definition 1. Range : In a network represented by G(V, E), range_i includes p_i and the set of its 1-hop neighbors. In this case, $|range_i|$ is equal to the degree of p_i in G plus 1. Note that ranges are symmetric i.e. $p_i \in range_i \Rightarrow p_i \in range_i$

Definition 2. Range Density : In a network represented by G(V, E), the range density, namely d, is equal to the size of the smallest range set of the network :

$$d \stackrel{def}{=} min(|range_i|), \forall p_i \in \Pi$$

We assume that d is known to every process.

Definition 3. f-Covering Network : A network represented by G(V, E) is f-covering if and only if G is (f + 1)-connected.

By Menger's Theorem[YG98], a graph G is (f+1)-connected if and only if it contains (f+1) independent paths between any two nodes. Thus, removing f nodes from G leaves at least *one* path between any pair of nodes (p_i, p_j) . Moreover, the range density d of the network will be greater than f+1, d>f+1. These lead to the following remark.

Remark 1. Let G(V, E) be an f-covering network, thus there is a path between any two nodes in G, in spite of f < n crashes.

4 Implementation of a Failure Detector of Class $\Diamond S$

This section presents a failure detector algorithm for a network where nodes do not move. The next section (5) extends this algorithm to support node mobility. This section firstly presents the principle of the *query-response* mechanism on which our algorithm is based. Then, it introduces some behavioral properties that, when satisfied by the underlying system, allow to implement a failure detector of the class $\Diamond S$. Based on such a properties, we propose an asynchronous failure detection algorithm. A proof that our implementation provides a failure detector of class $\Diamond S$ is also presented.

4.1 Query-Response Mechanism

The basic principle of our approach is the flooding of failure suspicion information over the network based on a local query-response mechanism. The algorithm proceeds execution by rounds. At each query-response round, a node broadcasts a QUERY message to the nodes of its range until it possibly crashes. The time between two consecutive queries is finite but arbitrary. A QUERY message sent by a node includes two sets of nodes: the set of nodes that it currently suspects of being faulty, and a set of the mistakes i.e., the nodes that were erroneously suspected of being faulty previously. Each node keeps a counter, which is incremented at every round. Every new information that is generated by this node about failure suspicions or correction of false suspicions (mistakes) within a round is tagged with the current value of such a counter. This tag mechanism avoids old information to be taken into account by nodes of the network.

Upon receiving a QUERY message from a node of its range, a node sends it back a RESPONSE message. A QUERY issued by a node is satisfied when it receives at least d-f corresponding RESPONSE messages. Moreover, each couple of QUERY- RESPONSE messages are uniquely identified in the system². Notice that we assume that a node issues a new QUERY only after the previous one is terminated. Moreover, when a node broadcasts a QUERY message, we assume that it receives the QUERY too, and that its own response always arrives among the first d-f responses it is waiting for.

 $^{^{2}}$ For the sake of simplicity, such identification is not included in the code of the algorithms of the paper.

4.2 Behavioral Properties

Let us define some behavioral properties that processes should have in order to ensure that our proposed implementation of a failure detector satisfies the properties of class $\Diamond S$ in an unknown network.

In order to implement any type of unreliable failure detector with an unknown membership, processes should interact with some others to be known. According to [FJA06], if there is some process in the system such that the rest of processes have no knowledge whatsoever of its identity, there is no algorithm that implements a failure detector with weak completeness, even if links are reliable and the system is synchronous. Thus, in order to implement a $\Diamond S$ failure detector, the following membership property, namely \mathcal{MP} , should be ensured by all processes in the system.

Property 1. Membership Property (\mathcal{MP}) . Let $t \in \mathcal{T}$. Denote $known_j^t$ the set of processes from which p_j has received a QUERY message at time t. Let K_i^t be the set of processes p_j that, at time t, have received a QUERY from p_i . That is, $K_i^t = \{p_j \mid p_i \in known_j^t\}$. A process p_i satisfies the membership property if:

$$\mathcal{MP}(p_i) \stackrel{def}{=} \exists t \geq 0 \in \mathcal{T} : |K_i^t| > f + 1$$

This property states that, to be part of the membership of the system, a process p_i (either correct or not) should interact at least once with other processes in its range by broadcasting a QUERY message. Moreover, this query should be received and represented in the state of at least one correct process in the system, beyond the process p_i itself.

Let us define another important property in order to implement a timer-free failure detector in a system with an unknown membership. It is the *responsiveness property*, namely \mathcal{RP} , which denotes the ability of a node to reply to a QUERY among the first nodes.

Property 2. Responsiveness Property (\mathcal{RP}). Let $t, u \in \mathcal{T}$. Denote $rec_from_j^t$ the set of d-f processes from which p_j has received responses to its QUERY message that terminated at or before t. The \mathcal{RP} property of the correct process p_i is defined as follows:

$$\mathcal{RP}(p_i) \stackrel{def}{=} \exists u \in \mathcal{T} : \forall t > u, \forall p_j \in range_i, \, p_i \in rec_from_j^t$$

Intuitively, the $\mathcal{RP}(p_i)$ property states that after a finite time u, the set of the d-f responses received by any neighbor of p_i to its last QUERY always includes a response from p_i .

4.3 Implementation of a Failure Detector of Class $\Diamond S$ for Unknown Networks

Algorithm 1 describes our protocol for implementing a failure detector of class $\Diamond S$ when the underlying system is an *f-covering* network, satisfying the behavioral properties.

We use the following notations:

- $counter_i$: denotes the round counter of node p_i .

- $suspected_i$: denotes the current set of processes suspected of being faulty by p_i . Each element of this set is a tuple of the form $\langle id, counter \rangle$, where id is the identifier of the suspected node and counter is the value of $counter_i$ when p_i generated the information that it suspected node id of being faulty.
- $mistake_i$: denotes the set of nodes which were previously suspected of being faulty but such suspicions are currently considered to be false. Similar to the $suspected_i$ set, the $mistake_i$ is composed of tuples of the form $\langle id, counter \rangle$ i.e, counter indicates when the information that id is falsely suspected was generated.
- rec_from_i : denotes the set of nodes from which p_i has received responses to its last QUERY message.
- $known_i$: denotes the current knowledge of p_i about its neighborhood. $known_i$ is then the set of processes from which p_i has received a QUERY messages since the beginning of execution.
- $Add(set, \langle id, counter \rangle)$: is a function that includes $\langle id, counter \rangle$ in set. If an $\langle id, \rangle$ already exists in set, it is replaced by $\langle id, counter \rangle$.

The algorithm is composed of two tasks. Task T1 is made up of an infinite loop. At each round, a QUERY message is sent to all nodes of p_i 's range (line 6). This message includes the set of nodes that p_i currently suspects and the set of mistakes of which p_i is aware. Node p_i waits for at least d-f responses, which includes p_i 's own response (line 7). Then, p_i detects new suspicions (lines 9-15). p_i starts suspecting each non previously suspected node p_j that it knows ($p_j \in known_i$) but from which it does receive a RESPONSE to its last QUERY. If a previous mistake information related to this new suspected node exists in the mistake set $mistake_i$, it is removed from it (line 12) and the counter $counter_i$ is updated to a value greater than the mistake tag (line 11). The new suspicion information is then included in $suspected_i$ with a tag which is equal to the current value of $counter_i$ (line 14). Finally, at the end of task T1, $counter_i$ is incremented by one (line 16).

Task T2 allows a node to handle the reception of a QUERY message sent by another node of its range. A QUERY message contains the information about suspected nodes and mistakes kept by the sending node. However, based on the tag associated to each piece of information, the receiving node only takes into account the ones that are more recent than those it already knows.

The two loops of task T2 respectively handle the information received about suspected nodes (lines 21–31) and about mistaken nodes (lines 32–37). Thus, for each node p_x included in the suspected (respectively, mistake) set of the QUERY message, p_i includes the node p_x in its $suspected_i$ (respectively, $mistake_i$) set only if the following condition is satisfied: p_i received a more recent information about p_x status (failed or mistaken) than the ones it has in its $suspected_i$ and $mistake_i$ sets. A more recent information is characterized by the fact that p_x has never been suspected or false suspected by p_i or by the fact that its counter in the p_i sets is less than the new received $counter_x$ (see lines 22 and 33). In such a case, p_i also removes the node p_x from its $mistake_i$ (respectively, $suspected_i$) set (lines 28 and 35).

Furthermore, in the first loop, a new mistake is detected if the receiving node p_i is included in the suspected set of the QUERY message (line 23). Then, p_i adds itself in its local

Algorithm 1 Asynchronous Implementation of a Failure Detector

```
1: init :
 2: \ suspected_i \leftarrow \varnothing; mistake_i \leftarrow \varnothing; counter_i \leftarrow 0
 3: known_i \leftarrow \emptyset
 4: Task T1:
 5:
     loop
          broadcast QUERY(suspected_i, mistake_i)
 6:
          wait until RESPONSE received from at least (d-f) distinct processes
 7:
 8:
          rec\_from_i \leftarrow the set of distinct nodes from which p_i has received a response at line 7
          for all p_j \in known_i \setminus rec\_from_i \mid \langle p_j, - \rangle \not \in suspected_i do
 9:
10:
               if \langle p_j, counter \rangle \in mistake_i then
11:
                   counter_i = max(counter_i, counter + 1)
                   mistake_i = mistake_i \setminus \langle p_j, - \rangle
12:
13:
               end if
14:
               Add(suspected_i, \langle p_i, counter_i \rangle)
          end for
15:
16:
          counter_i = counter_i + 1
17: end loop
18: Task T2:
19: upon reception of QUERY (suspected_j, mistake_j) from p_j do
20: known_i \leftarrow known_i \cup \{p_j\}
21: for all \langle p_x, counter_x \rangle \in suspected_i do
          \mathbf{if}\ \langle p_x, - \rangle \not \in suspected_i \cup mistake_i\ or\ \langle p_x, counter \rangle \in suspected_i \cup mistake_i\ |\ counter < counter_x
22:
     then
23:
               if p_x = p_i then
24:
                   counter_i = max(counter_i, counter_x + 1)
25:
                    Add(mistake_i, \langle p_i, counter_i \rangle)
26:
               else
27:
                    Add(suspected_i, \langle p_x, counter_x \rangle)
28:
                   mistake_i = mistake_i \setminus \langle p_x, - \rangle
29:
               end if
30:
          end if
31: end for
     for all \langle p_x, counter_x \rangle \in mistake_j do
32:
33:
          \textbf{if } \langle p_x, - \rangle \not \in suspected_i \cup mistake_i \ or \ \langle p_x, counter \rangle \in suspected_i \cup mistake_i \ | \ counter \leq counter_x
     then
34:
               Add(mistake_i, \langle p_x, counter_x \rangle)
35:
               suspected_i = suspected_i \setminus \langle p_x, - \rangle
36:
          end if
37: end for
38: send RESPONSE to p_i
```

mistake set (line 25). The tag $counter_i$ associated to this mistake is equal to the maximum of the current value of $counter_i$ and the tag associated to the suspicion of p_i , included in $suspected_j$ set, incremented by one (line 24). At the end of task T2 (line 38), p_i sends to the querying node a RESPONSE message.

4.4 Example of the Execution of the Algorithm

Figure 1 illustrates an execution which shows the strong completeness property of Algorithm 1. We consider an 1-covering network (f = 1) whose range density is equal to 3. Thus, each querying node should wait for at least 2 responses (one from itself and the other from one of its neighbors).




Fig. 1 – Example of Failure Detection

We do not show a scenario from the beginning of execution of the algorithm, but one where every node i is already aware of the participants of its range $(known_i)$, see step (a). In step (b), A fails. Thus, as neither node B nor node C receive a responses from A to their respective QUERY, they start suspecting A. At the moment of the QUERY $counter_B$ is equal to 5 (see $suspected_B$) but $counter_C$ is equal to 10 (see $suspected_C$). Then, both B and C propagate their suspected sets to their neighbors in their next respective QUERY rounds as shown in step (c). Nodes D and E will include the corresponding information $\langle A, 5 \rangle$ and $\langle A, 10 \rangle$ in their respective sets $suspected_D$ and $suspected_E$. Node B will update its $suspected_B$ set since the counter of the received information from C is greater than the one that it keeps in its $suspected_B$. However, C will discard the information received from B. Similar to step (c), in step (d) nodes B, C, D and E include in their next QUERY message their respective suspected sets. Therefore, eventually the information $\langle A, 10 \rangle$ related to the failure of A will be delivered to all correct nodes of the network.

4.5 Proof

We present in this section a sketch of proof of both the strong completeness and eventual weak accuracy properties of our algorithm that characterize failure detectors of class $\Diamond S$ for an *f-covering network* composed of non-moving nodes.

Consider that the most recent status about a process p_x is stored in a supected or mistake set and represented by the tuple $\langle p_x, ct_x \rangle$ which has the greatest counter ct_x in the network. In case of equality between a suspicion and a mistake, we give arbitrarily precedence to the mistake.

Lemma 1. Consider an f-covering network. Let p_i be a correct process. Consider that, at time t, p_i owns the most recent status about p_x in the network $(\langle p_x, ct_x \rangle)$ in its suspected set (respectively, mistake_i set). If no more recent information about p_x status is generated afterward, then eventually all correct nodes will include $\langle p_x, ct_x \rangle$ in their suspected set (respectively, mistake set).

Démonstration. Since p_i is correct, it will execute line 6 and broadcast a QUERY message containing $\langle p_x, ct_x \rangle$ in the $suspected_i$ set (respectively, $mistake_i$ set) to all its neighbors. As channels are reliable, this QUERY message is received by every correct process $p_j \in range_i$. Thus, p_j will execute lines 21-31 (respectively, lines 32-37). Since ct_x is the greatest counter associated with p_x in the network, p_j executes line 27 (respectively, line 34) and add $\langle p_x, ct_x \rangle$ to its own $suspected_j$ set (respectively, $mistake_j$ set). In the next round, p_j , the same as p_i , must broadcast this new status regarding p_x in its respective sets. Thus, due to the f-covering network property, all nodes in the network eventually add $\langle p_x, ct_x \rangle$ in their suspected set (respectively, mistake set) and the lemma follows.

Lemma 2. Consider an f-covering network in which all processes satisfy \mathcal{MP} . Let p_f be a faulty process. If process p_i is correct then eventually p_f is permanently included in its suspected_i set.

Démonstration. Let us consider that p_f crashes at time t.

Remark 1. Since $\mathcal{MP}(p_f)$ is satisfied, p_f has sent to processes in $range_f$ at least one QUERY message before it crashed at time t. Then, a number of correct processes within $range_f$ will include p_f in their respective known set which is updated when a process receives a QUERY (line 20). Let us denote K this set of processes. Notice that, by \mathcal{MP} , |K| > f + 1, and then there is at least one correct process p_i such that $p_f \in known_i$.

Remark 2. As p_f has crashed, there will be a time t' > t after which all processes in K will never receive a RESPONSE message from p_f (i.e., $p_f \notin rec_from$ sets of processes within K) (line 7). Thus, if p_f was not already suspected by these processes (line 9), it will be included in their corresponding suspected sets with a tag equal to the current value of their respective counter or with a greater tag then the one associated with p_f in the mistake set if it was previously in there (line 14). At this point no more information about p_f can be generated since only p_f can generate a mistake about itself (line 23) and only processes in K can generate a new suspicion and p_f is already in their suspected set. Thus, the most recent information about p_f sent in a QUERY message is either (1) a suspicion

or (2) a mistake. In the first case, following Lemma 1, all correct processes will eventually include p_f in their respective suspected set. Since no new information about p_f is generated, p_f is permanently suspected by all correct nodes. In the second case, by Lemma 1, the mistake eventually reach a correct process p_i in K, which removes p_f from $suspected_i$. At the next round, p_i will include p_f in $suspected_i$ with a greater tag since $p_f \notin rec_from_i$ and $p_f \notin suspected_i$. This information will in turn be propagated to all correct processes, following the propagation Lemma 1. Thus, all correct processes will permanently suspect p_f since no new information about p_f is generated.

Lemma 3. Consider an f-covering network in which all processes satisfy \mathcal{MP} . Let p_i be a correct process which satisfies the responsiveness property $\mathcal{RP}(p_i)$. There is a time u after which p_i is not included in the suspected, set of any correct process p_i .

Démonstration. Remark 1. According to $\mathcal{RP}(p_i)$, there is a time t after which every process p_j in the neighborhood of p_i receives a RESPONSE message from p_i in reply to their query. Thus, after time t, p_i is always included in the rec_from sets of all nodes within its $range_i$. Since a process starts being suspected only if its reply is not received by one of its neighbor (lines 9-15), no process adds p_i to its suspected set due to a QUERY message sent after time t.

Remark 2. If p_i is not included in any suspected set in the network, clearly p_i cannot be suspected anymore. If p_i is included in at least one suspected set, there are two cases to consider: the most recent piece of information about p_i is either (1) a mistake or (2) a suspicion. In the first case, based on Lemma 1, all processes which were suspecting p_i will eventually execute lines 34-35 upon receiving the propagated mistake and remove p_i from their suspected set definitely. In the second case, following Lemma 1, p_i will eventually deliver a QUERY message with p_i in the suspected set. This will cause p_i to generate a new mistake with a greater tag. This mistake will in turn be propagated to all processes, which will remove p_i from their suspected set if they were suspecting it.

Theorem 1. Algorithm 1 implements a failure detector of class $\Diamond S$, assuming an f-covering network of non-moving nodes which satisfies the behavioral properties \mathcal{RP} , \mathcal{MP} and with f < n.

Démonstration. Consider a correct process p_i and a fault process p_f . To satisfy the strong completeness property, we must prove that eventually p_f is permanently included in $suspected_i$ set of p_i . This claim follows directly from Lemma 2. To satisfy the eventual weak accuracy property, we must prove that there is a time u after which p_i is not included in the $suspected_j$ set of any correct process p_j . This claim follows directly from Lemma 3 and the theorem follows.

5 Extension for Mobility Management

In this section we present an extension for Algorithm 1 that supports mobility of nodes. For such an extension some new behavioral properties in respect to mobility of nodes and the underlying system must be defined.

5.1 Mobility Behavioral Properties

Let p_m be a moving node. Notice that a node can keep continuously moving and reconnecting, or eventually crashes. Nonetheless, we consider that p_m should stay connected to the network for a sufficient period of time in order to be able to update its state with recent information regarding failure suspicions and mistakes. Otherwise, it would not update its state properly and thus completeness and accuracy properties of the failure detector would not be ensured. Hence, in order to capture this notion of "sufficient time of reconnection", the following mobility property, namely $\mathcal{M}obi\mathcal{P}$, has been defined.

Property 3. Mobility Property ($Mobi\mathcal{P}$). Let $t \in \mathcal{T}$. Let Q_i^t be the set of processes from which p_i has received a QUERY message that terminated before or at t. A process p_i satisfies the mobility property if:

$$\mathcal{M}obi\mathcal{P}(p_i) \stackrel{def}{=} \exists t \geq 0 \in \mathcal{T} : |Q_i^t| > f + 1$$

This property should be satisfied by all moving nodes when they reconnect to the network. Thus, $\mathcal{M}obi\mathcal{P}(p_m)$ ensures that, after reconnecting, there will be a time at which process p_m should have received QUERY messages from at least one correct process, beyond itself. Since QUERY messages carry the state of suspicions and mistakes in the membership, this ensures that process p_m will update its state with recent informations.

We assume also that the membership property holds for all moving nodes when they reconnect to the network. Thus, $\mathcal{MP}(p_m)$ ensures that, after reconnecting, there will be a time at which process p_m interacts at least once with other processes in its $range_m$, broadcasting a QUERY message which will be delivered by at least one correct processes in $range_m$, beyond p_m .

Regarding the underlying system behavior, we consider that despite mobility, the f-covering property of the network is ensured and that the range density d of the network does not change. Moreover, we have extended the \mathcal{RP} property such that neighbors of a node p, which has the \mathcal{RP} property, eventually stop moving outside p's range. Otherwise, even if p has the \mathcal{RP} property, a moving node would add p in its known set whenever it belonged to p's range and then it would suspect p when it moved outside p's range. The extension of \mathcal{RP} property, namely $\mathcal{M}obi\mathcal{RP}$, is defined as follows:

Property 4. Mobility Responsiveness Property ($Mobi\mathcal{RP}$). Let $t \in \mathcal{T}$. Denote $range_i^t$ the set of processes in $range_i$ at t. A process p_i satisfies the mobility responsiveness property if:

$$\mathcal{M}obi\mathcal{RP}(p_i) \stackrel{def}{=} \mathcal{RP}(p_i) : \exists u \in \mathcal{T} : \forall t > u, \forall t' > t, p_j \in range_i^t \Rightarrow p_j \in range_i^{t'}$$

 $\mathcal{M}obi\mathcal{RP}$ should hold for at least one correct non-moving node.

5.2 Implementation of a Failure Detector of Class $\Diamond S$ for Mobile Unknown Networks

The extension of the algorithm to support mobility of nodes is based on the same query-response principle presented in Section 4. When a node p_m moves to another range, it starts being suspected of having crashed by those nodes of its old range, since it cannot reply to QUERY messages from the latter anymore. Hence, QUERY messages that include p_m as a suspected node will be propagated to nodes of the network. Eventually, when p_m reconnects to the network, it will receive such suspicion messages. Upon receiving them, p_m will correct such a mistake by including itself (p_m) in the mistake set of its corresponding QUERY messages. Such information will be propagated over the network. On the other hand, p_m will start suspecting the nodes of its old range since they are in its known set. It then will broadcast this suspected information in its next QUERY message. Eventually, this information will be corrected by the nodes of its old range, and the corresponding generated mistakes will spread over the network, following the same principle. Notice that, in order to avoid a "ping-pong" effect between information about failure suspicions and corrections (mistakes), a mechanism should be added to the algorithm in order to remove from known sets those nodes that belong to remote ranges.

In Algorithm 2, we just show the lines which need to be included in task T2 of Algorithm 1 in order to support mobility of nodes. Lines 36–38 should be added in the **if** block of the second loop of task T2, just after line 35 of Algorithm 1. They allow the updating of the known sets of both the moving node p_m and of those nodes that belong to the original range of p_m . For each mistake $\langle p_x, counter_x \rangle$ received from a node p_j such that node p_i keeps an old information about p_x , p_i verifies whether p_x is the sending node p_j . In they are different, p_x should belong to a remote range $range_x$, such that $p_x \notin range_i$. Thus, process p_x is removed from the local set $known_i$.

Algorithm 2 Asynchronous Implementation of a Failure Detector with Mobility of Nodes

```
36: if (p_x \neq p_j) then
37: known_i = known_i \setminus \{p_x\}
38: end if
```

5.3 Proof

We present in this section a sketch of proof of both the strong completeness and eventual weak accuracy properties of the extended algorithm 2 that characterize failure detectors of class $\Diamond S$ for an *f-covering network* composed of moving and non-moving nodes.

Lemma 4. (1) Infinitely often, during a run, the known_i set contains either correct processes which are in range_i or faulty processes. Moreover, (2) for every process p_i which satisfies $\mathcal{MP}(p_i)$, then there is a correct process p_i , such that $p_i \in known_i$.

Démonstration. Let us observe that the QUERY-RESPONSE messages are exchanged between processes in the same range. Thus, on the execution of line 20, the set $known_i$ is updated when p_i receives QUERY messages from other processes in its $range_i$. Beyond line 20, $known_i$ may be updated at lines 36–38, in order to remove nodes suspected to be in another range, different from p_i 's range. This may happen due to a mobility. Thus, if a process which raised a mistake (p_x) is different from the process who carries it (p_j) , probably p_x does not belong to $range_i$, because otherwise, p_i would have received the mistake by p_x itself. It may happen that p_x was in $range_i$ at some point in time, but due to a move, it has changed to another neighborhood, such that $p_x \not\in range_i$. Wherever the case, process p_i is going to remove p_x from its $known_i$ set and the part (1) of this lemma follows.

Let us prove part (2) of the lemma. Since $\mathcal{MP}(p_i)$ is satisfied, there is at least one correct process p_k which has received a QUERY message from p_i after p_i has connected or reconnected to the network at time t. Thus, $p_i \in known_k$. Nonetheless, later, p_i can be removed from $known_k$ by the execution of lines 36–38 due to a suspicion of mobility. But, notice that, since channels are reliable, the QUERY from p_i in which $p_i \in mistake_i$ is going to eventually arrive to p_k . In this case, two situations can occur. Situation (1). If this QUERY is the first one to arrive at p_k , it will satisfy the predicate of line 33, thus lines 34–35 are executed, but not lines 36-38. Afterward, when a QUERY from a process p_i arrives containing the mistake over p_i , and such that $p_i \neq p_j$, then since this mistake has already been taken into account, the predicate of line 33 will not be satisfied and lines 36-38 are not executed. Thus p_k will not remove p_i from $known_k$ set. Situation (2). A QUERY from a process p_i is the first one to arrive at p_k containing the mistake over p_i , and such that $p_i \neq p_j$. In this case, the predicate of line 33 is satisfied and lines 36-38 are executed. Thus p_k removes p_i from $known_k$. Nonetheless, later, a QUERY from p_i arrives in which $p_i \in mistake_i$. In this case, process p_k will execute line 20 including p_i in $known_k$. Moreover, since this mistake has already been taken into account, the predicate of line 33 will not be satisfied and lines 36–38 are not executed. Thus p_k will not remove p_i from $known_k$ set. This concludes the proof of part (2).

Lemma 5. Consider an f-covering network in which all nodes satisfy \mathcal{MP} and all moving nodes satisfy $\mathcal{M}obi\mathcal{P}$. Lemma 1 holds for every correct process p_i (moving or non-moving).

Démonstration. The lemma follows directly from Lemma 1 if p_i is a non-moving node. To take into account moving nodes, we should consider two cases. Case (1). Assume that p_i is a correct moving node which has the most recent status about process p_x . As soon as p_i reconnects to the network at time t', it will execute line 6 and broadcast a QUERY message to all its neighbors. Since $\mathcal{MP}(p_i)$ holds, p_i is correct and channels are reliable, every correct node $p_j \in range_i$ receives this QUERY message. Since, $|range_i| > f+1$, there will be at least one correct non-moving node p_k which receives this QUERY. Thus, by the same arguments of Lemma 1, the lemma follows.

Case (2). Assume that p_i is a correct moving node which has not yet the most recent status about process p_x and let us consider that due to Lemma 1, every non-moving node

has added p_x in its suspected (respectively, mistake) set before or at time t. As soon as p_i reconnects to the network at time $t' \geq t$, since $\mathcal{M}obi\mathcal{P}(p_i)$ is satisfied, p_i will receive QUERY messages from at least a correct process p_j with the last status of suspicion and mistaken informations about p_x . Thus, p_i will eventually add p_x in its suspected_i (respectively, $mistake_i$) set and the lemma follows.

Lemma 6. Consider an f-covering network in which all nodes satisfy \mathcal{MP} and all moving nodes satisfy \mathcal{MobiP} . Let p_f be a faulty process (moving or non-moving). If process p_i (moving or non-moving) is correct then eventually p_f is permanently included in its suspected set

Démonstration. If p_i and p_f are non-moving nodes, the lemma follows directly from Lemma 2. To take into account moving nodes, let us assume that p_i is a correct moving node which has the most recent status about process p_f . Due to Lemma 4 and the same arguments of Lemma 2 (Remark 1), p_f is in the known set of at least one correct process in the network. We should consider the following cases.

Case (1). Consider that p_f crashes at time r < t. Let us suppose that p_i is the only correct process such that $p_f \in known_i$. Moreover, before broadcasting this information to its neighborhood, p_i moves at time t. Since p_i keeps its state during the moving, $p_f \in suspected_i$. When p_i reconnects to the network at time t', due to Lemma 5, this information about the suspicion of p_f will be propagated to all correct nodes in the network. Finally, due to the same arguments of Lemma 2 (Remark 2) and Lemma 5, p_f is permanently included in every suspected set of a correct process, either moving or non-moving.

Case (2). Consider that p_f crashes at time s, t < s < t'. Suppose that p_i has p_f in its $mistake_i$ when it starts moving at time t. Since p_i keeps its state during the moving, $p_f \in mistake_i$ when p_i reconnects to the network at time t'. Since p_i has the most recent status about p_f , then, due to Lemma 5, this information about the mistake of p_f will be propagated to all correct nodes in the network. Nonetheless, as soon as p_f is faulty, due to the same arguments of Lemma 2 and Lemma 5, p_f is permanently included in every suspected set of a correct process, either moving or non-moving.

Lemma 7. Consider an f-covering network in which all nodes satisfy \mathcal{MP} and all moving nodes satisfy \mathcal{MobiP} . Let p_i be a correct non-moving node which satisfies the mobility responsiveness property $\mathcal{MobiRP}(p_i)$. There is a time u after which p_i is not included in the suspected_j set of any correct process p_j (moving or non-moving).

Démonstration. Since $Mobi\mathcal{RP}(p_i)$ is satisfied, there is a time s after which, $\mathcal{RP}(p_i)$ holds and nodes in the neighborhood of p_i do not leave $range_i$. Thus, due Lemma 3 (Remark 1), there is a time s' after which, no process in the network adds p_i to its suspected set (on to the execution of lines 9–15).

Due to Lemma 3 (Remark 2), we can ensure that p_i will not be included in any suspected set of non-moving correct nodes. We must then prove that eventually p_i is not included in

the $suspected_m$ set of any correct moving node p_m . Let us consider a correct moving node p_m starting to move at time t and stopping to move at time t'. Notice that, if p_m does not suspect p_i before moving at time t, the claim follows from Lemma 3 (Remark 2). Suppose that p_m suspects p_i before or at time t. Then, since p_m keeps its state during the moving, $p_i \in suspected_m$ when p_m reconnects to the network at time t'. If the suspicion over p_i represents the most recent information in the network, due to Lemma 5, it is going to be diffused to all correct nodes. Nonetheless, as soon as p_i is correct, p_i will revoke such a suspicion by the execution of lines 23-26, which will generate a new mistake with a greater tag. Due to Lemma 5, this mistake will be propagated to all correct processes, then p_m will permanently remove p_i from its $suspected_m$ set.

Theorem 2. Algorithm 2 implements a failure detector of class $\Diamond S$, assuming an f-covering network of moving and non-moving nodes which satisfies the behavioral properties \mathcal{RP} , \mathcal{MP} , $\mathcal{M}obi\mathcal{P}$ and $\mathcal{M}obi\mathcal{RP}$.

Démonstration. The strong completeness property follows directly from Lemma 6. The eventual weak accuracy property follows directly from Lemma 7 and the theorem follows.

6 Performance Evaluation

In this section we study and evaluate the behavior of our asynchronous failure detector compared to a timer-based one. To this end, we have chosen the gossip-based heartbeat unreliable failure detector proposed by Friedman and Tcharny in [FT05].

Our performance experiments were conducted on top of the OMNeT++ discrete event simulator [omn]. We assume a two-dimensional region S of 700mx700m. Transmission range r is set to 100m in all runs. The number of nodes N is fixed to 100 and each simulation lasts 30 minutes. The one-hop network delay δ is equal to 1ms in average. Since our unreliable failure detector needs a network where the $f_covering$ property always holds, the N nodes can not be placed randomly inside the region S. The initial topology of the network is in fact gradually built before the beginning of execution of an experiment. Thus, we start by inserting a graph clique of f+2 nodes organized in a circle whose radius is equal to r/2. Then, at each step, a new node of S is randomly chosen. The latter is included in the network regardless it has f+1 neighbors in the current configuration. The construction of the network stops when it reaches N nodes.

In the unreliable FD proposed by Friedman and Tcharny, a node periodically sends heartbeat messages to its neighbors. A vector is included in every heartbeat message such that each entry in the vector corresponds to the highest heartbeat known to be sent from the corresponding node. Every Δ time units, each node increments the entry of the vector corresponding to itself and then broadcasts its heartbeat to its neighbors. Based on the performance experiments described in the authors's article, we have set Δ to 1s. Upon receiving a heartbeat message, a node updates its vector to the maximum of its local vector and the one included in the message. A node also associates a timer to each other node of

the system. Thus, node j set the timer of i to Θ whenever it receives a new information about i. On the other hand, if the timeout of i expires, it is considered suspected by j. Note that the value of Θ should take into count higher communication delay due to longer paths between two nodes. We have set the value of Θ to 2s.

Concerning the implementation of our FD, it is not feasible that a node continuously broadcasts a QUERY message since the network would be overloaded with messages. To overcome this problem, we have included a delay of Δ units of time between lines 7 and 8 of the Algorithm 1. Similar to the Friedman and Tcharny's approach, we have set Δ to 1s. However, by adding this waiting period, a processes may receive more than d-f replies. Therefore, the extra replies will also be included in the rec_from set of this process (line 8), reducing then the number of false suspicions. It is worth remarking that this improvement does not change the protocol correctness.

6.1 Failure Detection

In order to evaluate the completeness property of both failure detectors, we have measured the impact of the range density d of the network on their respective failure detection time (Figure 2). The number of faults is equal to 5 and they are uniformly inserted during an experiment. The range density d varies from 7 to N/2 nodes. For each density, we have measure the average, maximum and minimum failure detection time.




Fig. 2 – Failure detection time vs. density

We observe that for both failure detectors there is no false suspicion. Furthermore, the propagation of failure suspicions is quite fast because the diameter of the network is relatively small. In the case of Friedman and Tcharny's FD, the mean failure detection time is always between $\Theta - \Delta$ and Θ time units, independently of d since failures are detected based on heartbeat vector values and timers. Such limit values can be explained: if node i crashes just after node j has set its timer related to i to Θ , j will detect the crash of i after Θ units

of time; if i crashes just before broadcasting a heartbeat, i.e. just after Δ units of time, j will detected the crash of i after $\Theta - \Delta$ units of time. On the other hand, for our FD, the failure detection time decreases with the range density. This happens because failure detection information is included in QUERY messages which spreads faster over the network when the density increases. We can notice that for values of d greater than 22, the failure detection time is uniform and equals around $\Delta + \delta$.

The maximum failure detection time characterizes the time for all nodes to detect a failure (strong completeness). We can observe that compared to Friedman and Tcharny's FD, this time is smaller and homogeneous for our FD, which can be also explained by the above mentioned propagation of failure information in QUERY messages.

6.2 Impact of mobility

We have evaluated the accuracy property when a node m which has 7 neighbors and is located at one boundary of the network moves about 500m at a speed of 2m/s. It starts moving at time 100s. We consider that while moving, node m does not interact with the other nodes as if it travels through a disturbance region where it can not send or receive any message. Thus, m stops executing while it moves. Furthermore, all neighbors of m must have d-f+1 neighbors. Such restriction is necessary to guarantee that at least d-f nodes will reply to the query of these old neighbors of m after it moves. The range density d of the network is equal to 7 and there is no fault.




Fig. 3 – Total number of false suspicions

For each experiment, the total number of false suspicions has been measured. Figure 3 shows the moment just before and after node m stops moving at time 356s. We can observe that all N-1 nodes suspect m before this time in both failure detectors. After it, false suspicions about node m start being corrected by all nodes. In Friedman and Tcharny's FD, there are no more false suspicions in around 1.5s. False suspicions about node m will also

start being corrected in our FD since m generates a mistake which is propagated over the network. However, node m at the same time starts suspecting its 7 old neighbors. Thus, it broadcasts such suspicions in its next QUERY message. This information spreads over the network and nodes of the system will start suspecting them too. This is the reason why the total number of false suspicions starts increasing after 357s till 358s when almost all nodes suspect the 7 old neighbors of m. However, at this time such an information also reaches the latter that then generate the corresponding mistakes and broadcast them. Such mistakes are propagated to all nodes of the network. All false suspicions are corrected by all nodes at 359.5s.

7 Related Work

As in our approach, some scalable failure detector implementations do not require a fully connected network. Larrea et al. proposed in [LFA00] an implementation of an unreliable failure detector based on a logical ring configuration of processes. Thus, the number of messages is linear, but the time for propagating failure information is quite high. In [GCG01], Gupta et al. proposed a randomized distributed failure detector algorithm which balances the network communication load. Each process randomly chooses some processes whose aliveness is checked. Practically, the randomization makes the definition of timeout values difficult. In [BMS03], a scalable hierarchical failure adapted for Grid configurations is proposed. However, the global configuration of the network is initially known by all nodes. It is worth remarking that none of these works tolerate mobility of nodes.

Few implementations of unreliable failure detector found in the literature focus on MA-NET environments. All of them are timer-based ones. In the Friedman and Tcharny algorithm [FT05], authors assumes a known number of nodes and that failures include message omissions too. In [TTS04], the authors exploit a cluster-based communication architecture for implementing a failure detector service able to support message losses and node failures. However, they provide probabilistic guarantees for the accuracy and completeness properties.

Sridhar presents in [Sri06] the design of a hierarchical failure detection which consists of two independent layers: a local one that builds a suspected list of crashed neighbors of the corresponding node and a second one that detects mobility of nodes across network, which corrects possible mistakes. Contrarly to our approach that allows the implementation of FD of class $\Diamond S$, the author's failure detector is an eventually perfect local failure detector of class $\Diamond P$ i.e., it provides strong completeness and eventual strong accuracy but with regard to a node's neighborhood.

In order to solve the problem of reaching agreement in mobile networks where processes can crash, Cavin et al. [CSS05] have adapted the failure detector definition of [CT96] to the case where the participants are unknown. They have introduced the concept of local participant detectors, which are oracles that inform the subset of processes that participating in the consensus. The authors construct an algorithm that solves consensus with an unknown number of participants in a fail-free network. Furthermore, they extend their solution and prove that a perfect failure detector (\mathcal{P}) is required for solving the fault-tolerant consensus

with a minimum degree of connectivity. Greve et al. [GT07] have subsequently extended this work, by providing a solution for the consensus in a fail-prone network which considers the minimal synchrony assumption (i.e., the $\Diamond S$), but at the expenses of requiring a higher degree of connectivity involving with the set of participants. We believe that our proposed $\Diamond S$ FD will be of great interest to implement this consensus algorithm over a MANET.

8 Conclusion

This paper has presented a new implementation of an unreliable failure detector for dynamic networks such as MANETs, where the number of nodes is not initially known and the network is not fully connected. Our algorithm is based on a query-response mechanism which is not timer-based. We assume that the network has the f-covering property, where f is the maximum number of failures. This property guarantees that there is always a path between two nodes despite of failures. Our algorithm can implement failure detectors of class $\Diamond S$ when both the behavioral responsiveness $(\mathcal{RP}, \mathcal{M}obi\mathcal{RP})$, membership (\mathcal{MP}) and mobility $(\mathcal{M}obi\mathcal{P})$ properties are satisfied by the underlying system. The proposed algorithm supports mobility of nodes as well. As a future work, we plan to adapt our algorithms and properties to implement other classes of failure detectors.

Références

- [BMS03] M. Bertier, O. Marin, and P. Sens. Performance analysis of a hierarchical failure detector. In *Proc. of the Int. Conf. on Dependable Systems and Networks*, San Francisco, CA, USA, june 2003. 3, 20
- [CSS05] D. Cavin, Y. Sasson, and A. Schiper. Reaching Agreement with Unknown Participants in Mobile Self-Organized Networks in Spite of Process Crashes. Technical report, 2005. 20
- [CT96] T. D. Chandra and S. Toueg. Unreliable failure detectors for reliable distributed systems. *Journal of the ACM*, 1996. 3, 4, 20
- [DT00] B. Devianov and S. Toueg. Failure detector service for dependable computing. In *Proc. of the First Int. Conf. on Dependable Systems and Networks*, pages 14–15, juin 2000. 3
- [FJA06] A. Fernández, E. Jiménez, and S. Arévalo. Minimal system conditions to implement unreliable failure detectors. In Proc. of the 12th Int. Symposium Pacific Rim Dependable Computing, pages 63–72. IEEE Computer Society, 2006. 7
- [FLP85] M. J. Fischer, N. A. Lynch, and M. S. Paterson. Impossibility of distributed consensus with one faulty process. *Journal of the ACM*, 32(2):374–382, apr 1985. 3
- [FT05] R. Friedman and G. Tcharny. Evaluating failure detection in mobile ad-hoc networks. *Int. Journal of Wireless and Mobile Computing*, 1(8), 2005. 3, 17, 20

[GCG01]	I. Gupta, T. D. Chandra, and G. S. Goldszmidt. On scalable and efficient dis-
	tributed failure detectors. In Proc. of the twentieth annual ACM symposium on
	Principles of distributed computing, pages 170–179. ACM Press, 2001. 3, 20

- [GT07] F. Greve and S. Tixeuil. Knowledge conectivity vs. synchrony requirements for fault-tolerant agreement in unknown networks. In Proc. of the Int. Conf. on Dependable Systems and Networks, Edinburgh, UK, June 2007. 21
- [LFA00] M. Larrea, A. Fernández, and S. Arévalo. Optimal implementation of the weakest failure detector for solving consensus. In *Proc. of the 19th Annual ACM Symposium on Principles of Distributed Computing*, pages 334–334, NY, July 16–19 2000. ACM Press. 3, 20
- [MMR03] A. Mostefaoui, E. Mourgaya, and M. Raynal. Asynchronous implementation of failure detectors. In *Proc. of Int. Conf. on Dependable Systems and Networks*, June 2003. 3
- [omn] OMNet++ Discret Event Simulation System. http://www.omnetpp.org. 17
- [SM01] I. Sotoma and E. Madeira. Adaptation algorithms to adaptative fault monitoring and their implementation on corba. In *Proc. of the IEEE 3rd Int. Symposium on Distributed Objects and Applications*, pages 219–228, september 2001. 3
- [Sri06] N. Sridhar. Decentralized local failure detection in dynamic distributed systems. The 25th IEEE Symposium on Reliable Distributed Systems, 0:143–154, 2006. 20
- [TTS04] A. Tai, K. Tso, and W. Sanders. Cluster-based failure detection service for large-scale ad hoc wireless network applications. In Proc. of the Int. Conf. on Dependable Systems and Networks, pages 805–814, New York City, USA, June 2004. IEEE Computer Society Press. 3, 20
- [YG98] Jay Yellen and Jonathan L. Gross. *Graph Theory & Its Applications*. CRC Press, 1998. 6

Table des matières

1	Introduction	i
2	2 Chandra-Toueg's Failure Detectors	
3 Model		5
4 Implementation of a Failure Detector of Class $\Diamond S$		6
	4.1 Query-Response Mechanism	(
	4.2 Behavioral Properties	7
	4.3 Implementation of a Failure Detector of Class $\Diamond S$ for Unknown Networks	7
	4.4 Example of the Execution of the Algorithm	10
	4.5 Proof	11

5	Extension for Mobility Management 5.1 Mobility Behavioral Properties	14
6	Performance Evaluation 6.1 Failure Detection	
7	Related Work	20
8	Conclusion	21



Unité de recherche INRIA Rocquencourt Domaine de Voluceau - Rocquencourt - BP 105 - 78153 Le Chesnay Cedex (France)

Unité de recherche INRIA Futurs : Parc Club Orsay Université - ZAC des Vignes
4, rue Jacques Monod - 91893 ORSAY Cedex (France)
Unité de recherche INRIA Lorraine : LORIA, Technopôle de Nancy-Brabois - Campus scientifique
615, rue du Jardin Botanique - BP 101 - 54602 Villers-lès-Nancy Cedex (France)
Unité de recherche INRIA Rennes : IRISA, Campus universitaire de Beaulieu - 35042 Rennes Cedex (France)
Unité de recherche INRIA Rhône-Alpes : 655, avenue de l'Europe - 38334 Montbonnot Saint-Ismier (France)
Unité de recherche INRIA Sophia Antipolis : 2004, route des Lucioles - BP 93 - 06902 Sophia Antipolis Cedex (France)





