A Delta Debugger for ILP Query Execution

Remko Tronçon* and Gerda Janssens

Katholieke Universiteit Leuven, Dept. of Computer Science,
Celestijnenlaan 200A, B-3001 Leuven, Belgium
{remko,gerda}@cs.kuleuven.be

Abstract. Because query execution is the most crucial part of Inductive Logic Programming (ILP) algorithms, a lot of effort is invested in developing faster execution mechanisms. These execution mechanisms typically have a low-level implementation, making them hard to debug. Moreover, other factors such as the complexity of the problems handled by ILP algorithms and size of the code base of ILP data mining systems make debugging at this level a very difficult job. In this work, we present the trace-based debugging approach currently used in the development of new execution mechanisms in hipP, the engine underlying the ACE Data Mining system. This debugger uses the delta debugging algorithm to automatically reduce the total time needed to expose bugs in ILP execution, thus making manual debugging step much lighter.

1   Introduction

Data mining [9] is the process of finding patterns that describe a large set of data best. Inductive Logic Programming (ILP) [12] is a multi-relational data mining approach, which uses the Logic Programming paradigm as its basis. ILP uses a generate-and-test approach, where in each iteration a large set of hypotheses (or 'queries') has to be evaluated on the data (also called 'examples'). Based on the results of this evaluation, the ILP process selects the "best" hypotheses and refines them further. Due to the size of the data of the problems handled by ILP, the underlying query evaluation engine (e.g. a Prolog system) is a crucial part of a real life ILP system. Hence, a lot of effort is invested in optimizing the engine to yield faster evaluation time through the use of new execution mechanisms, different internal data representations, etc.

   The development of new execution mechanisms for ILP happens mainly in the engine of the ILP system. These optimized execution strategies typically require a low level implementation to yield significant benefits. For example, the query pack [3] and adpack [17] execution mechanisms require the introduction of new dedicated WAM instructions, together with a set of new data structures which these instructions use and manipulate. Because of their low-level nature, finding bugs in the implementation of these execution mechanisms is very hard. While tracing bugs in these low-level implementations might still be feasible for

---

small test programs, many bugs only appear during the execution of the ILP algorithm on real life data sets. Several factors make debugging in this situation difficult:

- *The size of the ILP system itself.* Real life ILP systems group the implementation of many algorithms into one big system. These systems therefore often have a very large code base. For example, the ACE system [1] consists of over 150000 lines of code. In the case of the ACE system, the code base is very heterogeneous, where parts of code are written in different languages and others are generated automatically using preprocessors etc. This makes it in practice very hard to use standard tracing to detect bugs.
- *The complexity/size of the ILP problem.* With large datasets, it can take a very long time (hours, even days) before a specific bug occurs. When debugging, one typically performs multiple runs with small modifications to pin-point the exact problem, and so long execution times make this approach infeasible.
- *The high complexity of the hypothesis generation phase.* While the evaluation of hypotheses is often the bottleneck, some algorithms (such as rule learners) have a very expensive hypothesis generation phase. This phase is independent from the execution of the queries itself, and as such has no influence on the exposure of the bug. For algorithms with a very complex hypothesis generation, it can take a very long time for the bug in the execution mechanism to expose itself, even when the time spent on executing these queries is small.
- *Non-determinacy of ILP algorithms.* If an ILP algorithm makes random decisions (typically in the hypothesis generation phase), the exact point in time where the bug occurs changes from run to run. It is even possible that the bug does not occur at all in certain runs.

In [15], we proposed a trace-based approach for analyzing and debugging ILP data mining execution. This approach allowed easy and fast debugging of the underlying query execution engines, independent of the ILP algorithm causing the bug to appear. In this work, we present an extension to this debugging approach, automating a large part of the debugging process. By applying the *delta debugging algorithm* [18] on ILP execution traces, we automatically generate minimal traces exposing a bug, thus greatly reducing the time and effort needed to track the bug down. This approach is currently used in the development of new execution mechanisms in hipP [10], the engine underlying the ACE Data Mining system [1].

The organization of this paper is as follows: In Section 2, we give a brief introduction to Inductive Logic Programming. Section 3 discusses the collection of the run-time information necessary for our trace-based debugging approach. Section 4 then discusses applying the delta debugging algorithm on these traces to allow fast and easy debugging. We briefly discuss the implementation of our delta debugger in Section 5. Finally, we conclude in Section 6.

```
% QH: Queue of hypotheses
QH := Initialize
repeat
    Remove H from QH
    Choose refinements r_1, ..., r_k to apply on H
    Apply refinements r_1, ..., r_k on H to get H_1, ..., H_k
    Add H_1, ..., H_k to QH
    Evaluate QH
    Prune QH
until Stop-criterion(QH)
```

Fig. 1. Generic ILP Algorithm

2   Background: Inductive Logic Programming

The goal of Inductive Logic Programming is to find a theory that best explains a large set of data (or examples). More specifically, in the ILP setting at hand, each example is a logic program, and the theory is represented as a set of logical queries. Additionally, background knowledge about the problem domain is encoded as logical predicates.

ILP algorithms typically follow the same generate-and-test pattern: a set of queries is generated, of which all queries are tested on (a subset of) the examples in the data set. The query (or queries) which cover the data the best are then selected and extended, after which the process restarts with the extended queries. Hence, the query space is exhaustively searched, starting from the most general query and refining it further and further until the query (or queries) cover the data well enough. The generic ILP algorithm (as described in [12]) can be seen in Figure 1. In this algorithm, the `Initialize`, `Remove`, `Choose`, `Prune` and `Stop-criterion` procedures are to be filled in by the ILP algorithm, creating a special instance of this generic algorithm. Hence, these are the functions that characterize an ILP algorithm. In general, the most crucial step with respect to execution time is the *Evaluate* step: the (often large) set of queries $H_1, \ldots, H_n$ has to be run for each example. It is not uncommon to have a set of 3000 queries which are executed up to 1000 times. Therefore, fast query execution mechanisms are needed. Examples of these optimized techniques are query packs [3], adpacks [17], (lazy) control flow compilation [16], ... All of these techniques require a low-level implementation in the engine that the ILP algorithm uses. Due to the low-level nature of these optimized execution mechanisms, bugs in them are very hard to trace.

Concrete examples of ILP algorithms are TILDE [2], a decision tree learner, and WARMR [5], FOIL [13], and PROGOL [11], which are frequent pattern learners. Both algorithms were implemented in the ACE Data Mining system [1]. The ACE system uses hipP [10] as its execution engine, a high-performance Prolog engine (written in C) with specific extensions for ILP such as the above mentioned query optimization techniques. A typical ILP benchmarks is the Mutage-

```
query((atom(X,'c')), [1,2,3,4,5]).
query((atom(X,'h')), [1,2,3,4,5]).

query((atom(X,'c'),atom(Y,'o'), bond(X,Y)), [1,5]).
query((atom(X,'c'),atom(Y,'c'), bond(X,Y)), [1,5]).
```

Fig. 2. Example trace of an ILP run.

nesis data set [14], containing information about the structure of 230 molecules, and where the task of the ILP system is to learn to predict whether an unseen molecule can cause cancer or not. A more real-life data set is the HIV data set [6], containing over 4000 examples.

3 Gathering run-time information

Consider the generic ILP algorithm from Figure 1. The target of query execution optimizations is the *Evaluate* step, which takes a set of hypotheses to be evaluated, and evaluates them on the current dataset. The other steps that characterize the algorithm (such as finding suitable refinements for queries) are not important from an engine implementor's point of view. However, the latter are the most complex parts of the algorithm, and encompass most code of the algorithm itself. For our debugging purposes, we extract enough information from an ILP run necessary to be able to reproduce the *Evaluate* step, without running the ILP algorithm itself. More specifically, we only need to know the queries that the algorithm runs, and on which example each query is evaluated. How and why these queries were generated and selected is irrelevant for reconstructing the execution step.

To extract the desired information, we modify the *Evaluate* step from the ILP algorithm to record all evaluated queries to a file, which we call the *trace* of the algorithm. An example of such a trace after running a modified algorithm can be seen in Figure 2. This trace represents a run of an ILP algorithm that executed 4 queries: 2 queries that were executed on all 5 examples, and 2 extensions of the first query, which were only executed on the first and the last example. Notice that this trace is no longer dependent of the concrete algorithm itself, in the sense that it is just a sequence of queries the algorithm evaluated on the examples.

The gathered information can now be run through a trace simulator which, using the example database and background knowledge of the ILP algorithm, can now simulate the execution step of the ILP algorithm. Such a trace simulator is shown in Figure 3, and does nothing but run the original queries on the corresponding examples. While such a simulator in itself can be used for manually debugging query execution, we will also extend it further in Section 4 to yield an automatic debugging approach of different execution mechanisms.

```
% Run all queries from 'Trace' on 'Dataset'
simulate(Trace, Dataset) :-
    read(Trace, Input),
    ( Input == end_of_file ->
        true
    ;
        Input = query(Query, Examples),
        evaluate_query(Examples, Query, Dataset),
        simulate(Trace, Dataset)
    ).

% Evaluate a query on a set of examples
evaluate_query([], _, _).
evaluate_query([E|Es], Query, Dataset) :-
    load_example(Dataset, E),
    (call(Query), fail ; true),
    evaluate_query(Es, Query, Dataset).
```

Fig. 3. `simulate/2`: A simple trace simulator.

4   (Delta) debugging using traces

When developing optimizations for query evaluation, different execution mechanisms are investigated. When a new execution mechanism should yield the same final results as the existing ones, inconsistencies can be detected by running the ILP algorithm using each execution mechanism, and comparing the final results. For example, for TILDE, one can compare the learned decision trees to determine whether or not two runs are consistent with each other. This approach relies on the fact that new execution mechanisms speed up execution without changing the final results of the ILP algorithm. However, an inconsistent result only indicates that there is a bug in the execution *somewhere*, but it doesn't show *where*. To be able to determine this, the complete ILP algorithm has to be run using both the debugger of the host language of the ILP system (e.g. Prolog), and the debugger of the host language of the execution engine (e.g. C), where the actual bug of the execution mechanism is located. Because of the size and complexity of the ILP system, debugging on both levels simultaneously is very hard and time-consuming in practice. Moreover, testing execution mechanisms by comparing the output of the algorithm only works when the algorithm has deterministic behavior: if the decisions it makes are based on a random factor, the outputs of the algorithm can (slightly) differ, and comparing runs is not possible. This makes locating bugs in the implementation of optimizations even harder. Using execution traces for debugging solves many of these problems: trace execution is deterministic, and a trace simulator is so small that the focus of the debugging process is purely on the optimization itself. Moreover, traces can speedup debugging even more drastically by limiting execution to the part of the trace

causing the bug, as we show in this section.

When two runs of a deterministic ILP algorithm produce different results, this means that the query evaluation process selected different queries at some point. If the only difference between both runs is a query optimization scheme, this means that the optimization caused a query to succeed or fail where it did not before, meaning a bug (assuming that optimizations preserve success or failure of queries). Testing optimizations can therefore be reduced to comparing the success of query with and without the optimizations scheme. This can be achieved by simply running the trace through a simulator that records query successes, and runs every set of queries with and without the optimization enabled. Not only can such a simulator detect bugs this way, it can also pinpoint exactly in which query the bug occurs.

Due to the size of the trace, it might still be that a big part of the execution needs to be analyzed to find the bug. A bug occurring in a query is often also dependent on previously executed queries, which means that the trace cannot just be reduced to a single query to be able to reproduce and locate the bug. However, because the trace contains all the information determining the execution, locating a bug through traces can be turned into a *data slicing* [4] problem. The goal of data slicing is to take input data (i.e. a trace) that causes a bug to manifest itself, and reduce this data as much as possible to yield a smaller subset of data still exposing the bug. The standard approach to data slicing is simply to use binary search: split your data in two, test both halves, and continue with the half that still reproduces the bug. However, binary search might be too coarse-grained to find a bug, and as such fail to reduce the trace sufficiently. For example, if a bug occurs in the last query of the trace because of the execution of the first query, neither of both halves would reproduce the bug. *Delta Debugging* [18] is an automated data slicing technique that overcomes these issues. We describe delta debugging in the remainder of this section.

We briefly summarize the formalizations from [18]. Given a set of data $\mathcal{D}$ which causes a bug to appear. We denote this as $test(\mathcal{D}) = \text{fail}$. $\mathcal{D}_g \subseteq \mathcal{D}$ is a *global minimal data slice* if

$$test(\mathcal{D}_g) = \text{fail} \wedge \forall \mathcal{D}' \subseteq \mathcal{D} \cdot (|\mathcal{D}'| < |\mathcal{D}_g| \Rightarrow test(\mathcal{D}') \neq \text{fail})$$

In other words, $\mathcal{D}_g$ is the smallest possible subset of the original slice still reproducing the bug. Computing a global minimal data slice is infeasible in practice, since it requires testing of all $2^{|\mathcal{D}|}$ subsets of $\mathcal{D}$, which has exponential complexity. A less strict condition is the one of the *local minimum data slice* $\mathcal{D}_l$, for which no smaller subset exists that exposes the bug:

$$test(\mathcal{D}_l) = \text{fail} \wedge \forall \mathcal{D}' \subset \mathcal{D}_l \cdot test(\mathcal{D}') \neq \text{fail}$$

However, testing whether $\mathcal{D}_l$ is indeed a local minimum still requires $2^{|\mathcal{D}_l|}$ tests. An approximation to the local minimal slice is an *n-minimal data slice* $\mathcal{D}_n$,

function DDEBUG($\mathcal{D}$) :
    return DDEBUG($\mathcal{D}$,2)

function DDEBUG($\mathcal{D}$,n) :
    $\Delta_{i=1}^{n}$:= PARTITION($\mathcal{D}$,n)
    if $\exists \Delta_i, test(\Delta_i) = $ fail :
        return DDEBUG($\Delta_i, 2$)   – *'Reduce to subset'*
    else if $\exists \Delta_i, test(\mathcal{D} - \Delta_i) = $ fail :
        return DDEBUG($\mathcal{D} - \Delta_i, max(n-1, 2)$)   – *'Reduce to complement'*
    else if $n < |\mathcal{D}|$:
        return DDEBUG($\mathcal{D}, min(|\mathcal{D}|, 2n)$)   – *'Increase granularity'*
    else :
        return $\mathcal{D}$ – *'Done'*

Fig. 4. DDEBUG: The Delta Debugging algorithm. Finds a 1-minimal subset of $\mathcal{D}$ that causes the bug to appear.

which is a slice for which no n elements can be removed without making the bug disappear:

$$test(\mathcal{D}_n) = \text{fail} \wedge \forall \mathcal{D}' \subset \mathcal{D}_n \cdot (|\mathcal{D}_n| - |\mathcal{D}'| \leq n \Rightarrow test(\mathcal{D}') \neq \text{fail})$$

The delta debugging algorithm [18], depicted in Figure 4, finds a 1-minimal data slice of $\mathcal{D}$, i.e. a slice for which no one element can be removed without making the bug disappear. Note that even smaller slices might be constructed by removing more than one element. The algorithm works by dividing the data set in n (more or less) equal subsets, and checking if one of them still exposes the bug. If so, the process continues with this subset. If no subset exposes the bug but a complement of one of the subsets does, the process continues with the complement and increases granularity (such that the subsets in the next step are equally large). Otherwise, the granularity is increased if possible, or the process stops.

In our case, the data $\mathcal{D}$ corresponds to a trace, and every Δ_i represents a set of queries. Testing a Δ_i consists of running the trace with queries Δ_i through a trace simulator, and checking the output of the simulator for consistent results. For example, suppose that we have a query trace with 4 queries exposing a bug. Applying the delta debugging algorithm on the set of queries in the trace results in the steps from Figure 5. Note that some tests are repeated: a smart implementation can memorize tests, and re-use their answers. An important factor that determines the speed of the trace slicing is the granularity of the slicing process. Depending on what one considers the smallest part in which a trace can be divided, the delta debugger needs to consider more or less slices. A delta debugger for an ILP query trace can be set to use different granularities: it can either choose to find failing iterations in a trace, which gives fast results, but also less compact traces; it can prune the trace on the level of the queries themselves, giving a minimal trace; and, it can trim down the number of examples on which

Step	Call		Queries				Result
			1	2	3	4	
1	DDEBUG($\{1,2,3,4\}$,2)	Δ_1	●	●			√
		Δ_2			●	●	√
		Increase granularity					
2	DDEBUG($\{1,2,3,4\}$,4)	Δ_1	●				√
		Δ_2		●			√
		Δ_3			●		√
		Δ_4				●	√
		Δ_1^{-1}		●	●	●	×
		Reduce to complement					
3	DDEBUG($\{2,3,4\}$,3)						
		Δ_1		●			√
		Δ_2			●		√
		Δ_3				●	√
		Δ_1^{-1}			●	●	√
		Δ_2^{-1}		●		●	×
		Reduce to complement					
4	DDEBUG($\{2,4\}$,2)						
		Δ_1		●			√
		Δ_2				●	√
		Done: $\{2,4\}$ is 1-minimal					

Fig. 5. Example run of the delta debugging algorithm on a trace with 4 queries.

Fig. 6. Overview of the debugging process.

every query is run, reducing the number of times a query needs to be called to expose a bug. For example, consider the trace of Figure 2. If the delta debugger is set to find failing iterations, it only needs to perform two tests, one for every iteration. If it is set to find failing queries, it needs to consider each of the four queries separately, which introduces more checks than only finding the failing iteration. Finally, in the finest setting where every run of a query is trimmed down, the delta debugger needs to consider the combinations of the 14 runs (i.e. the first 2 queries are each run on 5 different example, whereas the last 2 queries are run on 2 different examples).

In the worst case, the DDEBUG algorithm needs to perform $|\mathcal{D}|^2 + 3|\mathcal{D}|$ tests. However, this worst case seldom occurs in practice. In the optimal case where there is only one element in the slice causing the bug to appear, the number of tests is bound by $2 \cdot log_2(|\mathcal{D}|)$.

5 Implementation

We have implemented and used the delta debugging approach in the development of new execution mechanisms for the ACE Data Mining system. An overview of the debugging process can be seen in Figure 6. The traces generated by the ILP algorithm are fed to the delta debugger, which trims it down to a smaller trace. The resulting trace is then fed into a trace simulator, and the engine (i.e. hipP) can then be manually debugged using the host language debugger (i.e. `gdb`).

We implemented two types of delta debuggers, which differ in the type of test they perform to detect when the execution of a trace exposes a bug. The simplest type of delta debugger is one that runs a trace through a trace simulator run in a separate hipP engine, and checks whether the process terminates successfully or not. This test can be used for bugs that cause an engine to fail (e.g. due to a segmentation fault). The second type of test compares the trace execution of two engines to check for inconsistent results. First, the queries from the original trace are adorned with extra goals, recording for every query in the

Trace	Granularity	Time	Tests	Resulting Trace		
				It[a]	Qu[b]	R[c]
1	Iterations	16.2s	10	1	137	822
	Queries	27.1s	26	1	1	6
	Queries ∘ Iterations	18.9s	24	1	1	6
	Examples ∘ Queries	27.6s	29	1	1	1
2	Iterations	78.0s	53	2	181	10942
	Queries	177.3s	157	2	2	236
	Queries ∘ Iterations	120.0s	136	2	2	236
	Examples ∘ Queries	180.4s	171	2	2	2
3	Iterations	138.1s	105	3	398	17235
	Queries	360.2s	338	3	3	265
	Queries ∘ Iterations	226.0s	271	3	3	265
	Examples ∘ Queries	371.1s	413	3	3	3

[a] Total number of iterations in the trace.
[b] Total number of queries in the trace.
[c] Total number of query runs necessary to reproduce the bug.

Table 1. Delta debugger execution time and number of tests performed for different granularities on three traces, together with statistics about the resulting traces. Traces are trimmed to the minimal amount of failing Iterations, Queries or Examples. Combinations of these granularities are denoted by ∘.

trace on which examples it succeeds. The test of the delta debugger then consists of calling hipP and running the resulting trace through both a plain trace simulator (see Figure 3) and a simulator with the (buggy) optimization enabled. The resulting logs of both runs are compared, and if they differ, the test fails. The delta debugger can be configured to use the different granularities described in Section 4: it can trim a trace to the minimal number of failing iterations, to the minimal number of failing queries, and, in its finest setting, to the minimal query runs (i.e. minimize both the number of queries and the examples they are run on).

Table 1 shows the execution time of the delta debugger using different combinations of granularities. For our experiments, we used a trace from a TILDE run on the Mutagenesis data set, with a lookahead setting of 2. The trace consists of 53 iterations of the algorithm, encompassing a total of 12908 queries. This trace was modified to get three variants: the first trace triggers an error when the last query of the last iteration is executed; the second trace triggers the same bug, yet only if the first query of the first iteration is executed as well; the third trace triggers the same bug whenever the first query and another query from the middle of the trace was executed. For each of these traces, the delta debugger was run using different granularities. Combinations of granularities are denoted by ∘, where $G_1 ∘ G_2$ means applying delta debugging with granularity G_1 on the trace resulting from delta debugging with granularity G_2. The delta debugger successfully minimized all three traces to the minimal trace needed to

reproduce the bug, being a trace of 1, 2 and 3 queries respectively. The results show that applying the delta debugging first on the level of iterations, and then pruning further on the query level requires less tests than immediately pruning the complete trace on the query level. Pruning on the iteration level gives a first 'rough' version of the trimmed down version of the trace, after which one can decide to prune further on the query level.

6   Conclusion

In this work, we presented a trace-based approach to debugging query execution mechanisms for ILP algorithms. Using traces to perform debugging yields several advantages. The specific workings of the ILP algorithm do not have to be known, as the traces are algorithm independent, yet provide enough information for performing a perfect simulation of the query execution of the algorithm itself. With trace-based execution, time is only spent on the execution of queries. Therefore, a complex query generation phase of an ILP algorithm does not affect the total execution time of a trace, and so debugging can be done faster. Finally, it is not necessary to have full knowledge of the code base of the ILP system, which can in practice become very large.

By applying the delta debugging algorithm on traces, the number of queries can be reduced significantly, allowing bugs to be exposed very fast without having to manually step through the complete trace.

In the past, traces of execution have been used to understand misbehavior of programs [7, 8]. These approaches do not use static traces, but instead interleave execution of the program with calls to the tracer, to avoid having to store the large traces. In the context of debugging ILP query execution, not storing the traces explicitly has the disadvantage that the execution times are higher (because time is spent in the ILP algorithm itself), and the bug might not occur if the algorithm is non-deterministic. Moreover, without a static trace, applying delta debugging to reduce the total time needed to expose a bug is not possible.

References

1. ACE. The ACE data mining system, 2006. http://www.cs.kuleuven.be/~dtai/ACE/.
2. H. Blockeel and L. De Raedt. Top-down induction of first order logical decision trees. *Artificial Intelligence*, 101(1-2):285–297, June 1998.
3. H. Blockeel, L. Dehaspe, B. Demoen, G. Janssens, J. Ramon, and H. Vandecasteele. Improving the efficiency of Inductive Logic Programming through the use of query packs. *Journal of Artificial Intelligence Research*, 16:135–166, 2002. http://www.cs.kuleuven.be/cgi-bin-dtai/publ_info.pl?id=36467.
4. T. W. Chan and A. Lakhotia. Debugging program failure exhibited by voluminous data. *Journal of Software Maintenance*, 10(2):111–150, 1998.
5. L. Dehaspe and H. Toivonen. Discovery of frequent datalog patterns. *Data Mining and Knowledge Discovery*, 3(1):7–36, 1999.

6. The developmental therapeutics program. U.S. Departement of Health and Human Services NIH, National Cancer Institute NCI. http://dtp.nci.nih.gov.

7. M. Ducassé. Coca: an automated debugger for c. In *ICSE '99: Proceedings of the 21st international conference on Software engineering*, pages 504–513. IEEE Computer Society Press, 1999.

8. M. Ducassé. Opium: An extendable trace analyser for Prolog. *The Journal of Logic programming*, 1999. Special issue on Synthesis, Transformation and Analysis of Logic Programs, A. Bossi and Y. Deville (eds), Also Rapport de recherche INRIA RR-3257 and Publication Interne IRISA PI-1127.

9. U. Fayyad and R. Uthurusamy. Data mining and knowledge discovery in databases. *Communications of the ACM*, 39(11), 1996.

10. hipP. hipP: A high performance Prolog system, 2006. http://www.cs.kuleuven.be/~dtai/hipp/.

11. S. Muggleton. Inverse entailment and Progol. *New Generation Computing, Special issue on Inductive Logic Programming*, 13(3-4):245–286, 1995.

12. S. Muggleton and L. D. Raedt. Inductive Logic Programming: Theory and methods. *Journal of Logic Programming*, 19/20:629–679, 1994.

13. J. Quinlan. Learning logical definitions from relations. 5:239–266, 1990.

14. A. Srinivasan, S. Muggleton, M. Sternberg, and R. King. Theories for mutagenicity: A study in first-order and feature-based induction. *Artificial Intelligence*, 85(1,2):277–299, 1996.

15. R. Tronçon and G. Janssens. Analyzing and debugging ILP data mining query execution. In *Proceedings of the Sixth International Workshop on Automated Debugging, AADEBUG 2005*, pages 105–109. ACM Press, 2005. http://www.cs.kuleuven.be/cgi-bin-dtai/publ_info.pl?id=41742.

16. R. Tronçon, G. Janssens, B. Demoen, and H. Vandecasteele. Fast frequent quering with lazy control flow compilation. *Theory and Practice of Logic Programming*, 2006. http://www.cs.kuleuven.be/cgi-bin-dtai/publ_info.pl?id=41995.

17. R. Tronçon, H. Vandecasteele, J. Struyf, B. Demoen, and G. Janssens. Query optimization: Combining query packs and the once-tranformation. In *Inductive Logic Programming, 13th International Conference, ILP 2003, Szeged, Hungary, Short Presentations*, pages 105–115, 2003. http://www.cs.kuleuven.be/cgi-bin-dtai/publ_info.pl?id=40938.

18. A. Zeller and R. Hildebrandt. Simplifying and isolating failure-inducing input. *Software Engineering*, 28(2):183–200, 2002. http://www.st.cs.uni-sb.de/papers/tse2002/.